Packet Switching on L2 (LAN Level)

Transparent Bridging (TB), Spanning Tree Protocol (STP), Rapid STP, L2 Bridging versus L3 Routing

Agenda

- Introduction
- Transparent Bridging Basics
 - Principles
 - Broadcast Storm Problem
- Spanning Tree Protocol
 - Introduction
 - Details
 - Convergence
- Rapid Spanning Tree Protocol
- Routing Principles

© 2008, D.I. Manfred Lindner

Transparent Bridging / STP, v4.12

Why Packet Switching on LAN?

LAN was primarily designed for shared media

- but too many clients cause performance problems
 - · higher probability for collisions in case of Ethernet

Partitioning of a network may be needed

- to achieve load balance
- for error limitation
- because of security reasons
- to deal with technological limitations
 - geographical expansion, number of clients
- for using resources of distant networks (remote networking)
- to join isolated LANs to corporate networks

© 2008, D.I. Manfred Lindner

Fransparent Bridging / STP, v4.12

Bridge / Router

network partitioning

- can be done by packet switches (OSI relay stations on L3)
 - already well known store and forward principle of WAN world

bridge

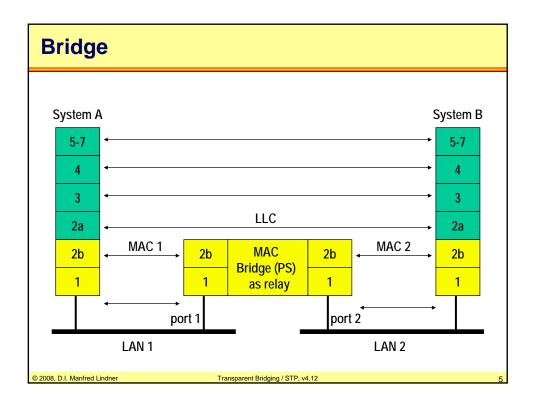
- packet switch implemented on OSI layer 2
- forwarding based on unstructured MAC addresses
- signposts stored in MAC bridging table (= routing table of L2 packet switch)
- above OSI Layer 2b bridges are transparent for all higher layer protocols (LLC and above)

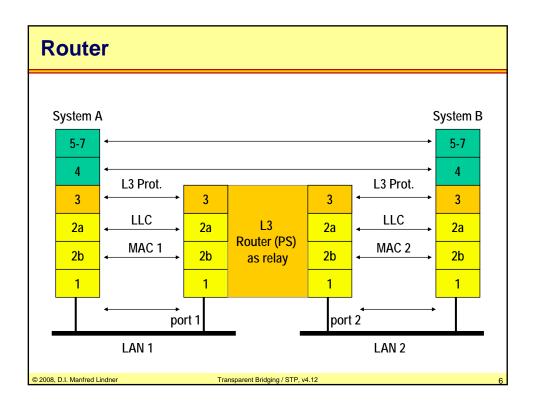
router

- packet switch implemented on OSI layer 3
- forwarding based on structured L3 addresses
- end system must speak the corresponding L3 "language"

© 2008, D.I. Manfred Lindner

Transparent Bridging / STP, v4.12





Bridging Methods

- Transparent Bridging (Ethernet world)
 - easy for end systems
 - more complex transferring-decisions (compared to source route bridging) have to be done by such a bridge
 - dedicated hardware
 - today's L2 switch technology (Ethernet switch) is actually a <u>fast</u> transparent bridge only
- Source Routing Bridging (Token Ring world, IBM)
 - more overhead for end systems (path finding)
 - source route bridges are less complex
 - could be done e.g. by a PC with two network cards
 - not further covered in this lecture

© 2008, D.I. Manfred Lindner

Transparent Bridging / STP, v4.12

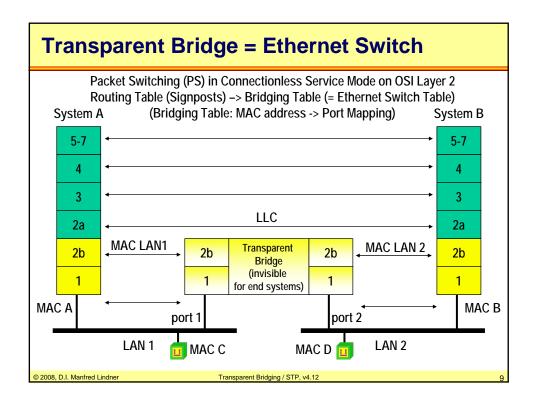
Agenda

- Introduction
- Transparent Bridging Basics
 - Principles
 - Broadcast Storm Problem
- Spanning Tree Protocol
 - Introduction
 - Details
 - Convergence
- Rapid Spanning Tree Protocol
- Routing Principles

© 2008, D.I. Manfred Lindner

Transparent Bridging / STP, v4.12

© 2008, D.I. Manfred Lindner



Transparent Bridging

1

- bridge is invisible for end systems
 - LAN 1 and LAN 2 appear to the end systems like one single, logical, big LAN -> transparent
- bridge uses layer 2 MAC-addresses
 - to decide if a given frame must be a forwarded or not
 - · destination-address of a frame is used for this
- MAC-addresses of all stations are registered in a bridging table
 - either statically done by administrator
 - or dynamically done by a self-learning mechanism
 - · source-address of a frame is used for this

© 2008, D.I. Manfred Lindner

Transparent Bridging / STP, v4.12

Transparent Bridging

2

• in case of a dynamic bridging table

- an aging mechanism allows for changes of MAC addresses in the network
 - may caused either by change of network card or by location change of end system
 - if an already registered MAC address is not seen within e.g. 5
 minutes as source address of a frame the bridging table entry is
 deleted

because of the transparency

 such a bridge must receive and process every frame on a LAN

flow control

originally not done between end systems and bridges

© 2008, D.I. Manfred Lindner

Transparent Bridging / STP, v4.12

..

Transparent Bridging

3

destination address of a frame is used for table look up which enables a simple decision:

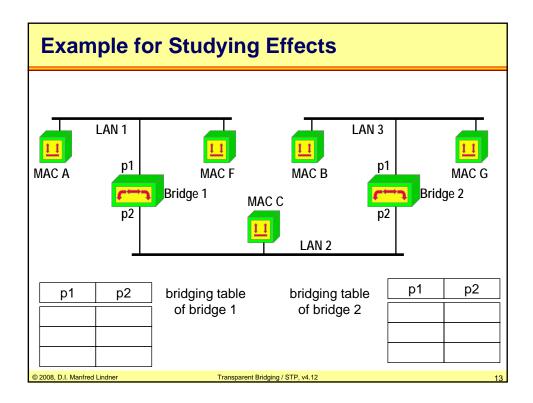
- <u>filtering</u>: frame will be rejected if destination's home is on the LAN segment of the receiving port
- forwarding: a duplicate of the frame will be forwarded to the appropriate port if destination's home is registered in the table of another port
- <u>flooding:</u> during learning time; frame will be forwarded to all other ports (multiport-bridge) if there is no entry in the table (unknown destination)

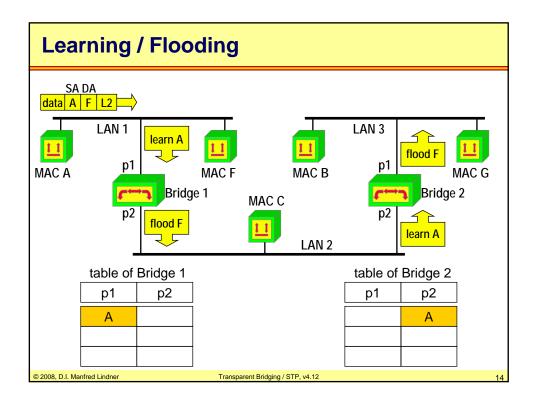
frames with broadcast/multicast-address

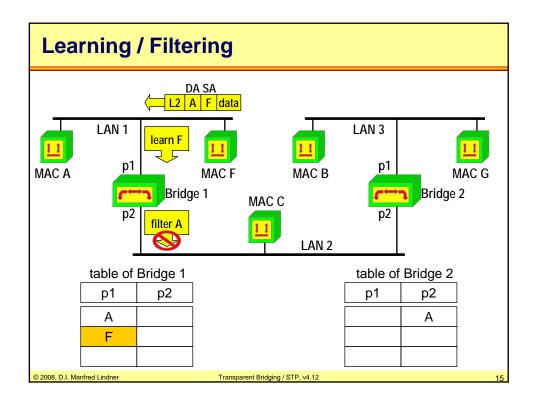
are always forwarded to all other ports

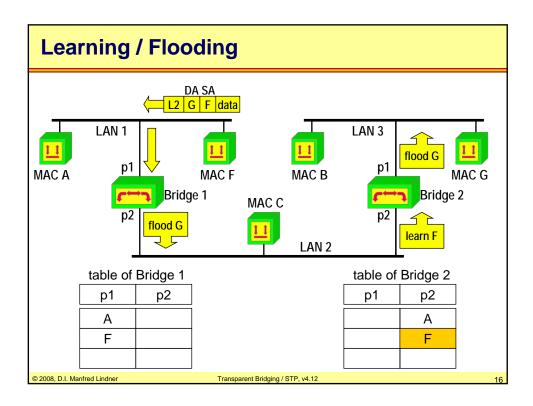
© 2008, D.I. Manfred Lindner

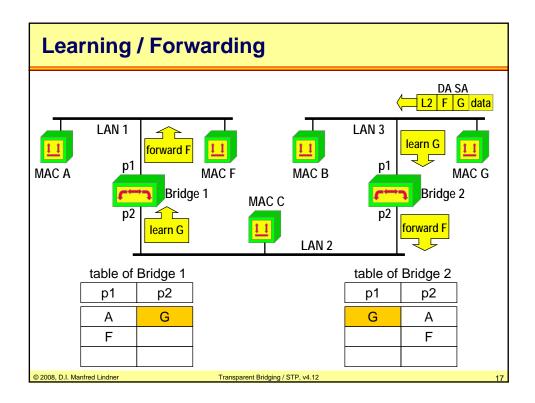
Transparent Bridging / STP, v4.12

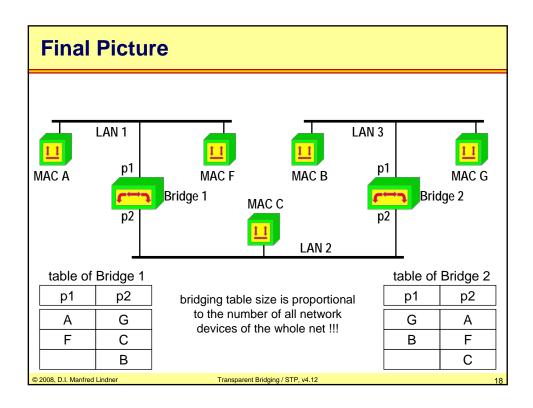


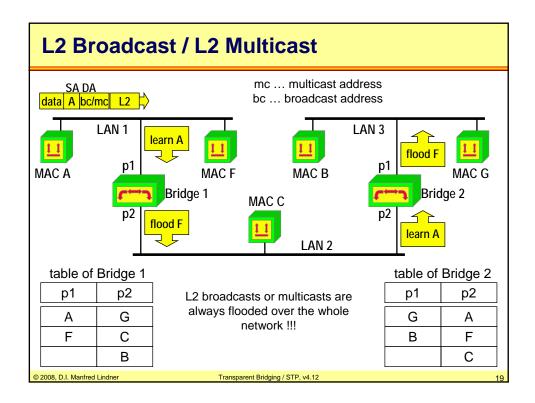












Transparent Bridging

- MAC addresses are unique but unstructured (U/L=0)
 - allows no determination of the location of a station
 - bridging table of all bridges will contain all MAC addresses of the whole net
- problem of location change
 - is solved by aging the table entries
- so network-stations
 - can keep their addresses on location changes
- note:
 - in Token Ring and Source Route Bridging environment MAC addresses will be structured (U/L=1)

© 2008, D.I. Manfred Lindner

Transparent Bridging / STP, v4.12

Collision Domain versus Broadcast Domain

bridge separates LAN into

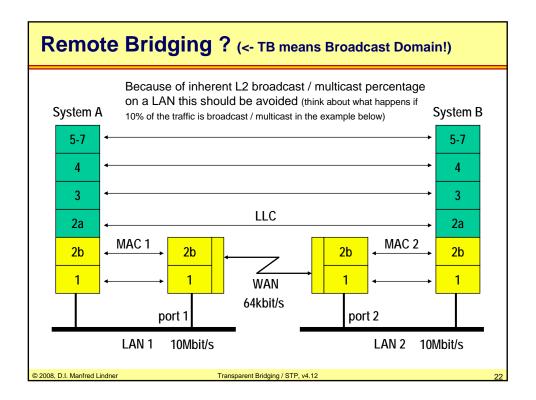
- multiple <u>collision domains</u>
- a collision on one LAN segment will not be seen on other LAN segments
- reason: store and forward behavior of bridge
- note: a repeater can't do that

but a bridged network is still

- one broadcast domain
- broadcast frames are always flooded
- reason: bridged LAN is still seen as one logical LAN by end systems, bridges are <u>transparent</u> for them

© 2008, D.I. Manfred Lindner

Transparent Bridging / STP, v4.12



Remote Bridging

- relay function of local bridge
 - is split in two half-bridges
- coupling of the half-bridges
 - via WAN-connection
- high amount of broadcast on a LAN
 - mismatch of data rates
 - slow WAN-connection can cause a buffer overflow in the bridge
 - therefore transparent bridging over WAN or any other
 Ethernet tunnelling technique should be avoided
- coupling of heterogeneous networks
 - not possible (e.g. one end system on LAN, other end system on WAN)

© 2008, D.I. Manfred Lindne

Transparent Bridging / STP, v4.12

Agenda

- Introduction
- Transparent Bridging Basics
 - Principles
 - Broadcast Storm Problem
- Spanning Tree Protocol
 - Introduction
 - Details
 - Convergence
- Rapid Spanning Tree Protocol
- Routing Principles

© 2008, D.I. Manfred Lindner

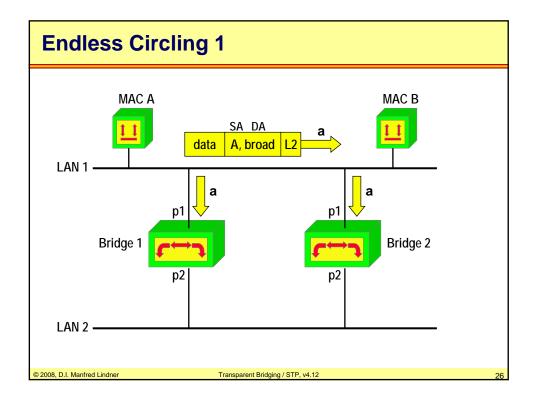
Transparent Bridging / STP, v4.12

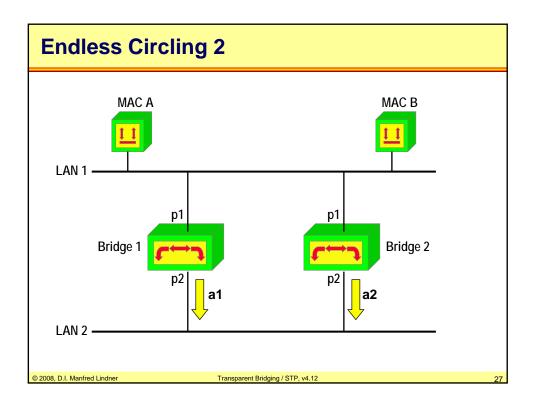
Parallel Paths

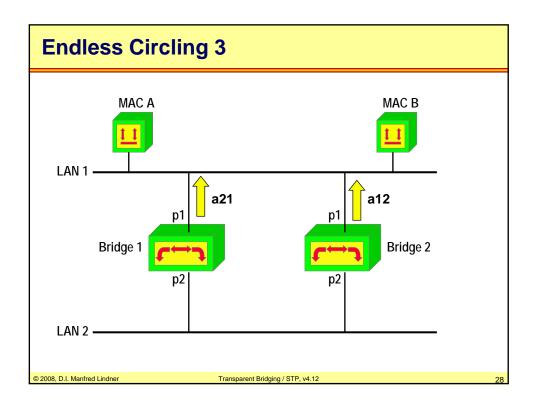
- parallel paths between two LAN segments cause
 - endless circling of frames with unknown destination
 - endless circling of broadcast-frames
 - endless circling of direct addressed frames during flooding phase
 - blocking of buffer-resources
- parallel paths in a more complex topology cause
 - overflow of all buffer-resources and stagnation of the LANs
 - Broadcast Storm
- to avoid these effects
 - Spanning Tree Protocol (STP)

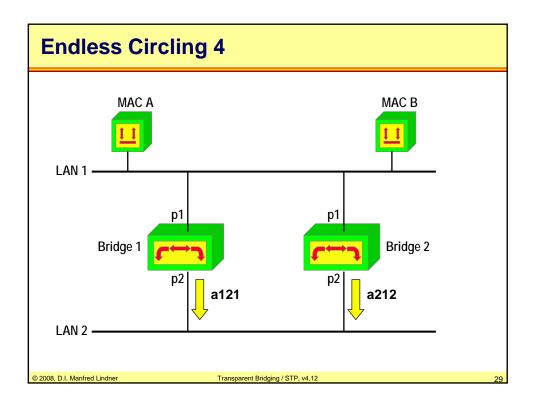
© 2008, D.I. Manfred Lindner

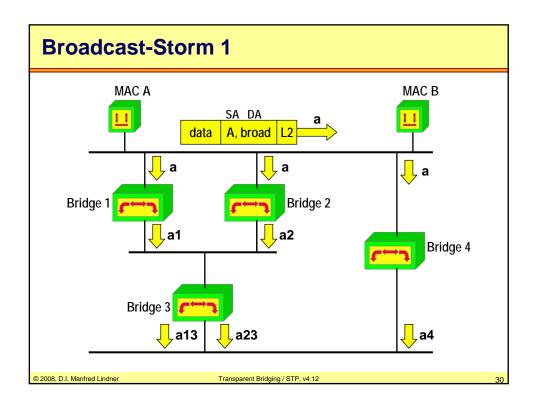
Transparent Bridging / STP, v4.12

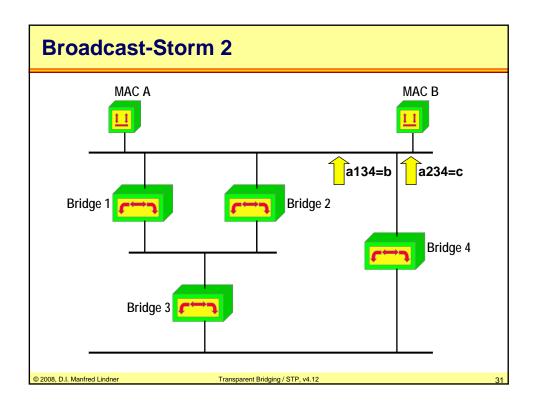


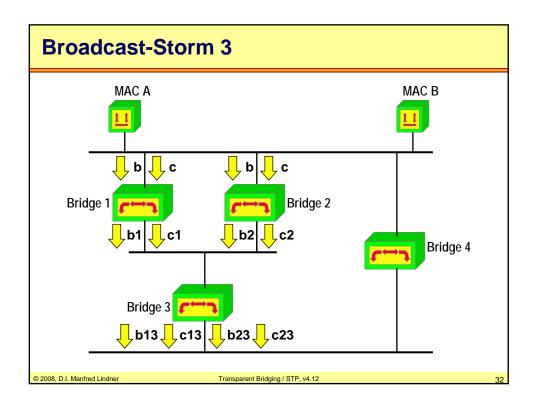


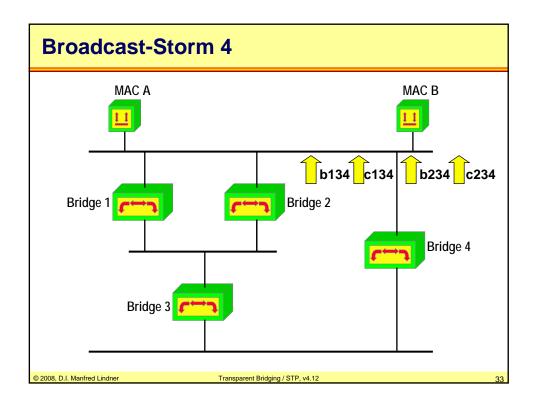


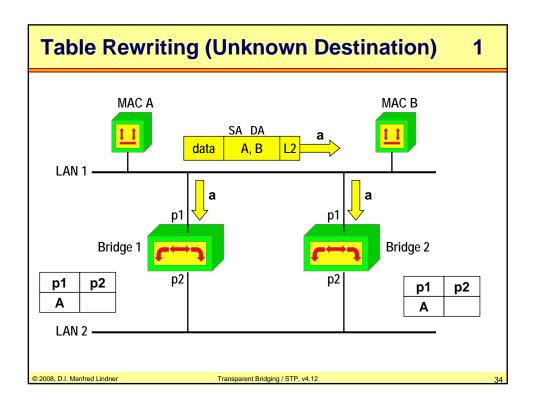


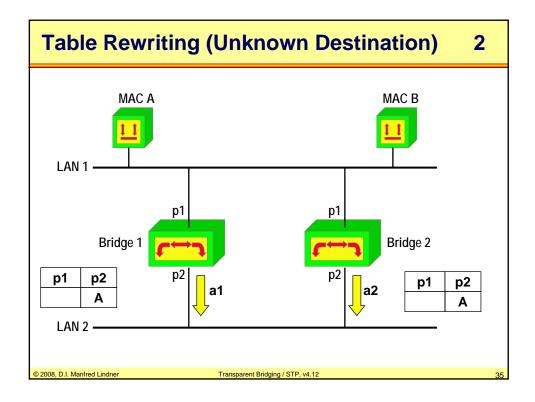


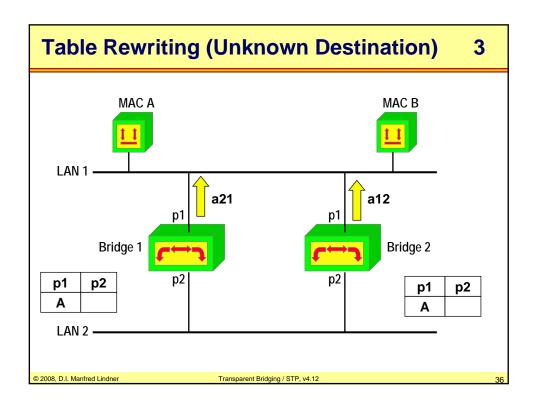












Agenda

- Introduction
- Transparent Bridging Basics
 - Principles
 - Broadcast Storm Problem
- Spanning Tree Protocol
 - Introduction
 - Details
 - Convergence
- Rapid Spanning Tree Protocol
- Routing Principles

© 2008, D.I. Manfred Lindner

Transparent Bridging / STP, v4.12

27

Spanning Tree (IEEE 802.1D)

- Spanning Tree Protocol (STP):
 - takes care that there is always exact only one active path between any 2 stations
 - implemented by a special communication protocol between the bridges
 - using BPDU (Bridge Protocol Data Unit) packets with MACmulticast address
 - failure of active path causes activation of a redundant path
- main disadvantage of STP
 - redundant lines or redundant network components cannot be used for load balancing

© 2008, D.I. Manfred Lindner

Transparent Bridging / STP, v4.12

Parameters for STP

1

Bridge Identifier (Bridge ID)

- combination of MAC-address and a priority number
 - typically, the lowest MAC-address of all ports is used for that
 - note: although bridge will not be seen by end systems, for bridge communication and management purposes a bridge will listen to one or more dedicated MAC addresses
 - Bridge-ID = priority# (2 Byte) + mac# (6 Byte)
- priority number can be configured by the administrator
 - default value is 32768
- lowest Bridge ID has highest priority
 - · lowest configured priority number
- if you keep default values
 - the bridge with the lowest MAC address will have the highest priority

© 2008, D.I. Manfred Lindner

Fransparent Bridging / STP, v4.12

Parameters for STP

2

Port Cost (C)

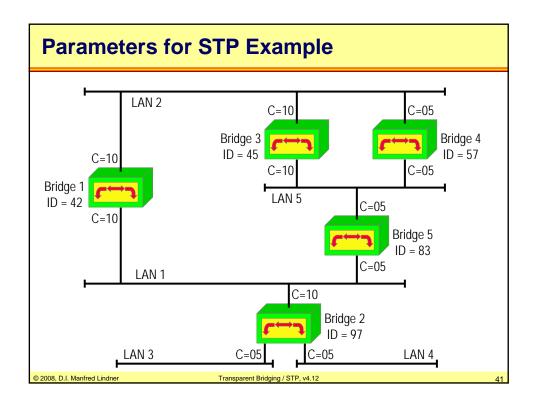
- costs in order to access local interface
- inverse proportional to the transmission rate
- default cost = 1000 / transmission rate in Mbit/s
 - so 10 Mbit/s Ethernet has a default Path Cost of 100
 - with occurrence of 1Gbit/s Ethernet rule was adapted
 100 Mbit/s = 19, 1Gbit/s = 4, 10Gbit/s = 2
- can be configured to a different value by the administrator

Port Identifier (Port ID)

- combination of port number and a priority number
 - Port-ID = port priority#.port#
- configured by the administrator
 - default port priority = 128

© 2008, D.I. Manfred Lindner

Transparent Bridging / STP, v4.12

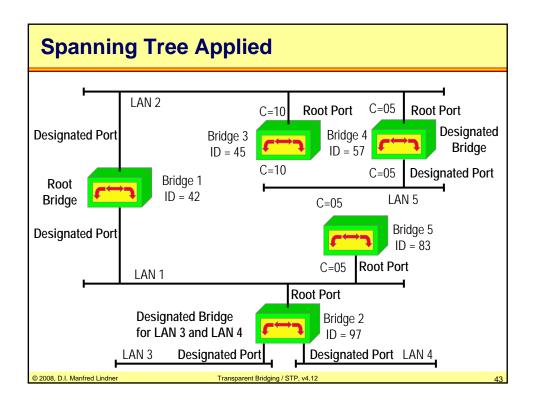


Spanning Tree Algorithm Summary

- select the root bridge
 - bridge with the lowest Bridge Identifier
- select the root ports
 - by computation of the shortest path from any other bridge to the root bridge
 - root port points to the shortest path towards the root
- select <u>one designated bridge</u> for every LAN segment which can be reached by more than one bridge
 - bridge with lowest root path costs on the root port side
 - corresponding port on other side is called designated port
- set the designated and root ports in <u>forwarding state</u>
- set all other ports in <u>blocking state</u>
- creates single paths from the root to all leaves (LAN segments) of the network

© 2008, D.I. Manfred Lindner

Transparent Bridging / STP, v4.12





Format of STP Messages - BPDU Format

Prot. ID		BPDU Type		Root ID	Root Path Costs	Bridge ID	Port ID	Mess. Age		Hello Time	
2 Byte	1 Byte	1 Byte	1 Byte	8 Byte	4 Byte	8 Byte	2 Byte	2 Byte	2 Byte	2 Byte	2 Byte

BPDU Bridge Protocol Data Unit (OSI term for this kind of message)

Root ID Who seems to be or who is the root bridge (R-ID)?

Root Path Cost How far is the root bridge away from me (RPC)?

Bridge ID ID of bridge transmitting this BPDU (O-ID)

Port ID port over which this BPDU was transmitted (P-ID)

© 2008, D.I. Manfred Lindner

ransparent Bridging / STP, v4.13

BPDU Fields

1

- Protocol Identifier:
 - 0000 (hex) for STP 802.1D
- Protocol Version:
 - 00 (hex) for version 802.1D (1998)
 - 02 (hex) for version 802.1D (2004)
- BPDU Type:
 - 00 (hex) for Configuration BPDU
 - 80 (hex) for Topology Change Notification (TCN) BPDU
- Root Identifier:
 - 2 bytes for priority (default 32768)
 - 6 bytes for MAC-address
- Root Path Costs in binary representation:
 - range 1-65535
- Bridge Identifier:
 - structure like Root Identifier

© 2008, D.I. Manfred Lindner

Transparent Bridging / STP, v4.12

BPDU Fields	2
--------------------	---

- Port Identifier:
 - 1 byte priority (default 128)
 - 1 byte port number
- Message Age (range 1-10s):
 - age of Configuration BPDU
 - transmitted by root-bridge initially using zero value, each passing-on (by designated bridge) increases this number
- Max Age (range 6-40s):
 - aging limit for information obtained from Configuration BPDU
 - basic parameter for detecting idle failures (e.g. root bridge = dead)
 - default 20 seconds
- Hello Time (range 1-10s):
 - time interval for generation of periodic Configuration BPDUs by root bridge
 - default 2 seconds

© 2008, D.I. Manfred Lindne

Fransparent Bridging / STP, v4.12

17

BPDU Fields

3

- Forward Delay (range 4-30s):
 - time delay for putting a port in the forwarding state
 - default 15 seconds
 - but that means 15 seconds listening plus 15 seconds learning
- Hello Time, Max Age, Forward Delay are specified by Root-Bridge
- Flags (a "1" indicates the function):
 - bit 8 ... Topology Change Acknowledgement (TCA)
 - bit 1 ... Topology Change (TC)
 - used in TCN BPDU's for signalling topology changes
 - TCN ... Topology Change Notification
 - in case of a topology change the MAC addresses should change quickly to another
 port of the corresponding bridging table (convergence) in order to avoid forwarding
 of frames to the wrong port/direction and not waiting for the natural timeout of the
 dynamic entry
 - the bridge recognizing the topology change sends a TCN BPDU on the root port as long as a CONF BPDU with TCA is received on its root port
 - bridge one hop closer to the root passes TCN BPDU on towards the root bridge and acknowledges locally to the initiating bridge by usage of CONF BPDU with TCA
 - when the root bridge is reached a flushing of all bridging table is triggered by the root bridge by usage of CONF BPDUs with TC and TCA set
 - the new location (port) is dynamically relearned by the actual user traffic

© 2008, D.I. Manfred Lindner

Transparent Bridging / STP, v4.12

MAC Addresses / LSAP / Network Diameter

• bridges use for STP-communication:

– multicast address:

0180 C200 0000 hex

0180 C200 0001 to 0180 C200 000F are reserved

0180 C200 0010 hex All LAN Bridges Management Group Address

- Note:
 - all addresses in Ethernet format
 - on the Token Ring the functional address: 0300 0000 8000 hex
- the L-SAP of LLC header

42 hex Bridge Spanning Tree Protocol

Maximum Bridge Diameter

 maximum number of bridges between any two end systems is 7 using default values for hello time, forward delay and max age

© 2008, D.I. Manfred Lindner

Transparent Bridging / STP, v4.12

40

Selection of Root Bridge

1

after power up

- all ports are set in a <u>Blocking State</u> and every bridge tries to become the <u>Root Bridge (RB)</u> of the Spanning Tree by sending Configuration BPDUs
- blocking state means:
 - end station Ethernet frames are not received and forwarded on such a port
 - but BDPU frames can still be received, manipulated by the bridge and transmitted on such port

© 2008, D.I. Manfred Lindner

Transparent Bridging / STP, v4.12

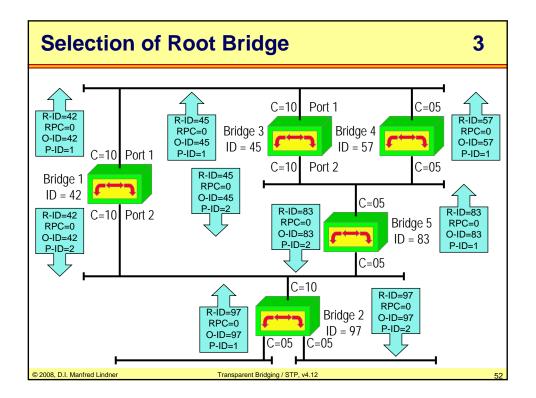
Selection of Root Bridge

2

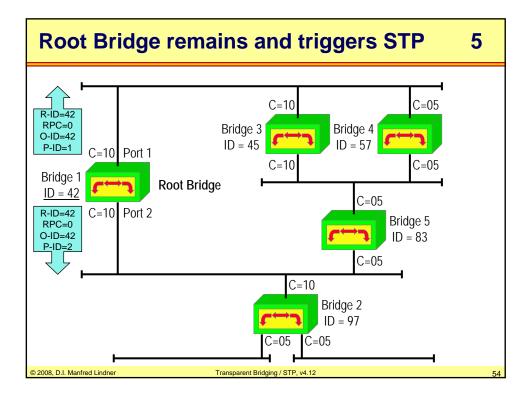
- using Configuration BPDUs, a bridge tells
 - which bridge actually is seen as RB
 - which path costs exist to this RB (Root Path Cost)
 - its own Bridge ID and Port ID
- bridge with the lowest Bridge ID becomes RB
- after selection of the RB
 - sending of Configuration BPDUs are exclusively triggered by the RB

© 2008, D.I. Manfred Lindner

ransparent Bridging / STP, v4.12



Selection of Root Bridge • strategy to determinate the RB: - if bridge receives a Configuration BPDU with lower Root Bridge ID as own Bridge ID • the bridge stops sending Configuration BPDUs on this port • the received and adapted Configuration BPDU is forwarded to all other ports - if bridge receives Configuration BPDU with higher Root Bridge ID as own Bridge ID • the bridge continues sending Configuration BPDUs with own Bridge ID as proposed Root Bridge ID on all ports • other bridges should give up



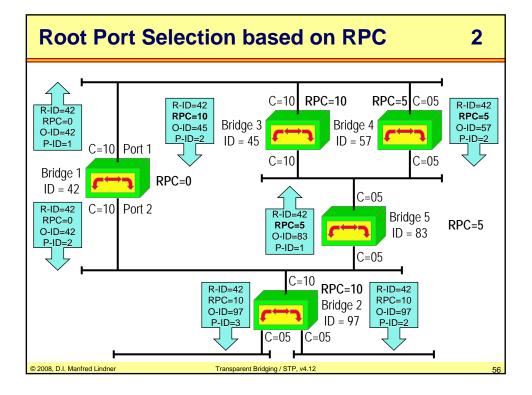
Root Port and Designated Bridge

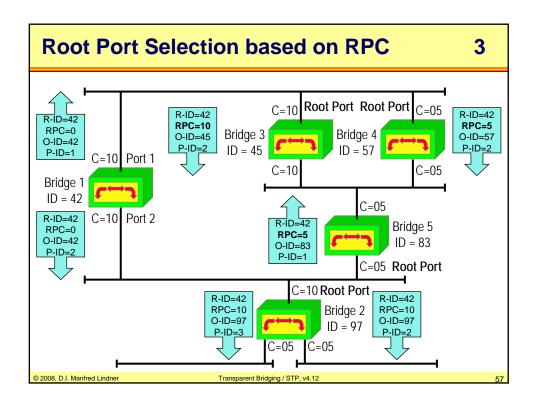
1

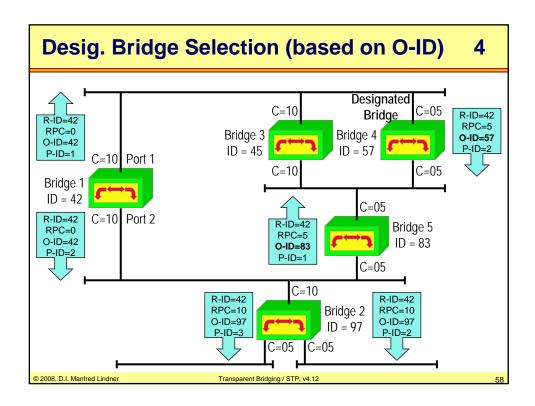
- now, every bridge determines which of its ports has the lowest Root Path Cost
 - Root Path Cost = sum of all port costs from this bridge to the RB, including port costs of all intermediate bridges
 - Calculation method: Root Path Cost received in BPDU + port cost of the local port receiving this BPDU
- this port becomes the Root Port
 - in case of equal costs the port ID decides (lower means better)
- similar to Root Bridge selection, a <u>Designated</u> <u>Bridge (DB)</u> is selected for each LAN-segment
 - bridge with the lowest Root Path Cost on its Root Port
 - in case of equal costs the bridge with the lowest Bridge ID wins

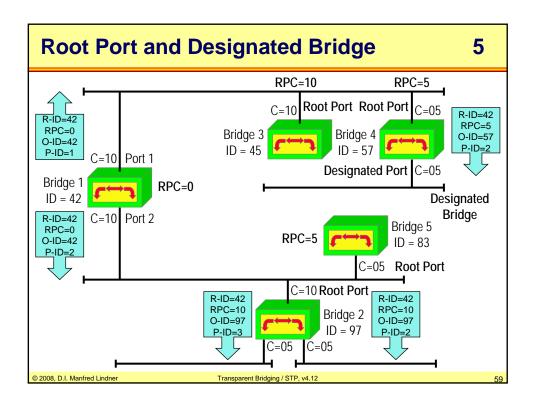
© 2008, D.I. Manfred Lindner

Transparent Bridging / STP, v4.12









Root Port and Designated Bridge

6

- using the Root Path Cost field in the Configuration BPDU, a bridge indicates its distance to the RB
- if a bridge receives a Configuration BPDU from a bridge which is closer to the RB
 - the receiving bridge adds its own port costs to the Configuration BPDU and forwards this message to all other ports
- if a bridge receives a Configuration BPDU from a bridge which is more distant to the RB
 - this receiving bridge drops this message and sends its own Configuration BPDU on this port
 - this Configuration BPDU contains its own Root Path Cost

© 2008, D.I. Manfred Lindner

Fransparent Bridging / STP, v4.12

Procedure Parameters Summary

Root Bridge

- lowest Bridge ID

Root Ports

- via Root Path Costs
 - sum of all port costs contained in the Configuration BPDU

Designated Bridge

- lowest Root Path Costs for a given LAN segment

© 2008, D.I. Manfred Lindner

Transparent Bridging / STP, v4.12

0.4

STP Configuration via Port States

every designated bridge

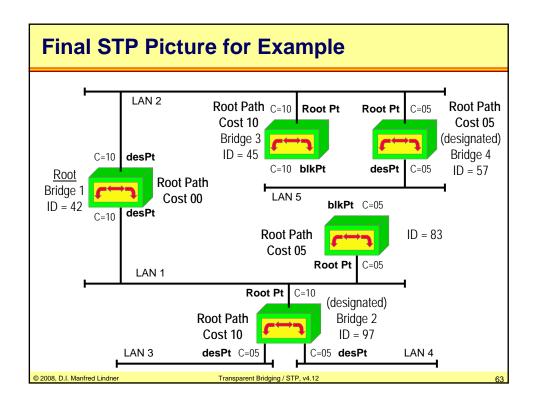
 declares its ports as designated ports and puts them (together with the Root Port) in the Forwarding State

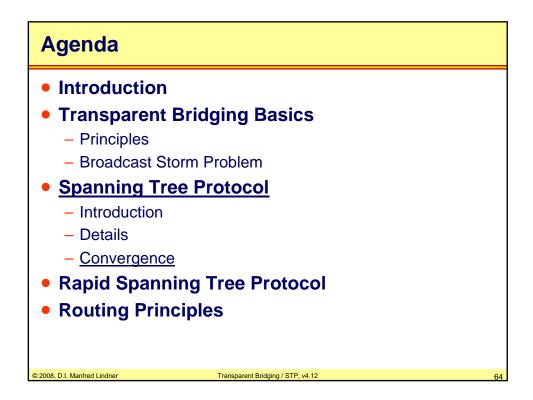
Root Bridge

- is automatically designated bridge for its LAN-segments and has only designated ports
- all other bridges (not RB or DBs)
 - keep their ports in the <u>Blocking State</u>
- from this moment on, the normal network operation is possible
 - there is only one path between any two arbitrary end systems

© 2008, D.I. Manfred Lindner

Transparent Bridging / STP, v4.1



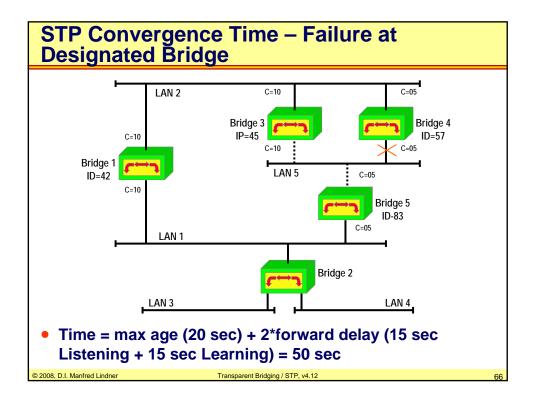


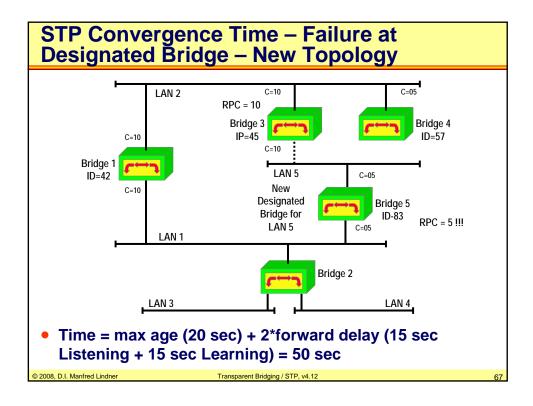
STP Error Detection

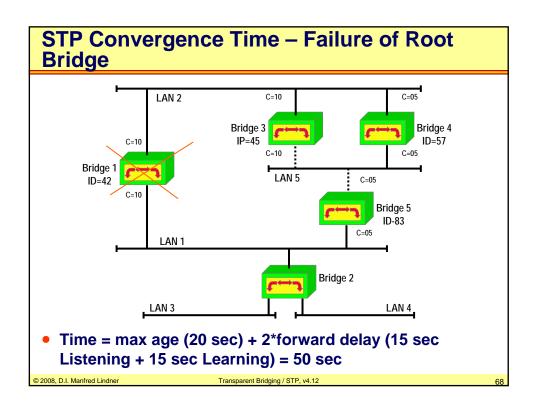
- normally the root bridge generates (triggers)
 - every 1-10 seconds (hello time interval) a Configuration BPDU to be received on the root port of every other bridge and carried on through the designated ports
 - bridges which are not designated are still listening to such messages on blocked ports
- if triggering ages out two scenarios are possible
 - root bridge failure
 - a new root bridge will be selected based on the lowest Bridge-ID and the whole spanning tree may be modified
 - designated bridge failure
 - if there is an other bridge which can support a LAN segment this bridge will become the new designated bridge

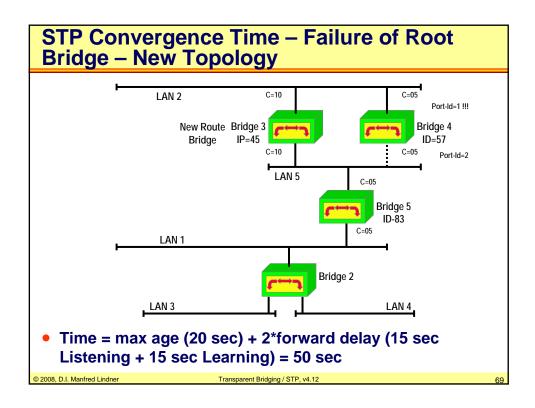
© 2008, D.I. Manfred Lindner

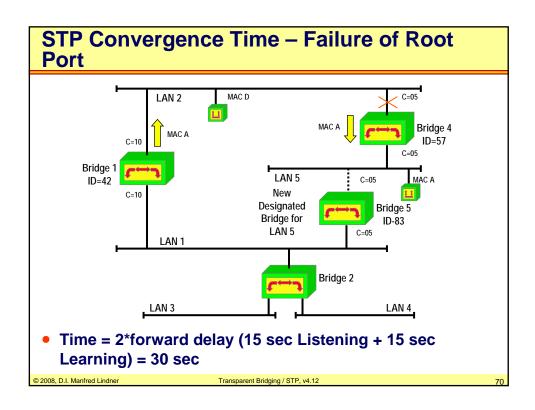
Fransparent Bridging / STP, v4.12

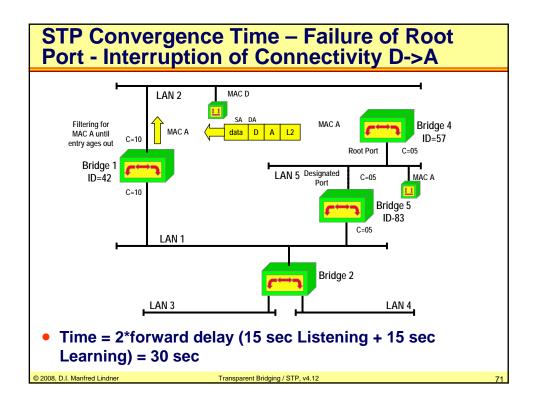


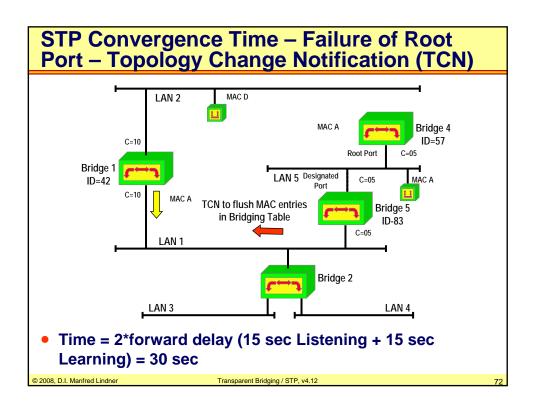












Some STP Facts

disadvantages of STP

- active paths are always calculated from the root, but the actual information flow of the network may use other paths
 - note: network-manager can control this via Bridge Priority, Path Costs und Port Priority to achieve a certain policy under normal operation
 - hence STP should be designed to overcome plug and play behaviour of default values
- redundant paths cannot be used for load balancing
 - redundant bridges may be never used if there is no failure of the currently active components
 - for remote bridging via WAN the same is true for redundant WAN links
- convergence time between 30 and 50 seconds
 - note: in order to improve convergence time Rapid Spanning Tree Protocol has been developed (802.1D version 2004)

© 2008, D.I. Manfred Lindner

Transparent Bridging / STP, v4.12

70

Agenda

- Introduction
- Transparent Bridging Basics
 - Principles
 - Broadcast Storm Problem
- Spanning Tree Protocol
 - Introduction
 - Details
 - Convergence
- Rapid Spanning Tree Protocol
- Routing Principles

© 2008, D.I. Manfred Lindner

Transparent Bridging / STP, v4.12

Introduction

Rapid Spanning Tree (RSTP)

- IEEE 802.1D version 2004 (former IEEE 802.1w)
- Can be seen as an evolution of the Spanning Tree Protocol (STP; IEEE 802.1D)
- Capable of reverting back to 802.1D version 1998
- Convergence time reduced to few seconds !!!

Terminology slightly changed

- Blocking port role is split into the Backup and Alternate port roles
 - Alternate port
 - Backup port
- Root port and Designated port roles still remain the same
- New port state
 - Discarding (see next slides for details)

© 2008, D.I. Manfred Lindner

Transparent Bridging / STP, v4.12

7-

Port States Comparison

STP (802.1d) Port State	RSTP (802.1w) Port State	Is Port included in active Topology?	Is Port learning MAC addresses?
disabled	discarding	No	No
blocking	discarding	No	No
listening	discarding	Yes	No
learning	learning	Yes	Yes
forwarding	forwarding	Yes	Yes

© 2008, D.I. Manfred Lindner

Fransparent Bridging / STP, v4.12

Port Roles

Root Port Role

 Receives the best BPDU (so it is closest to the root bridge)

Designated Port Role

- A port is designated if it can send the best BDPU on the segment to which it is connected
- On a given segment, there can be only one path toward the root-bridge

© 2008, D.I. Manfred Lindner

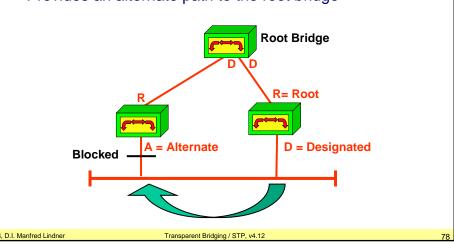
Transparent Bridging / STP, v4.12

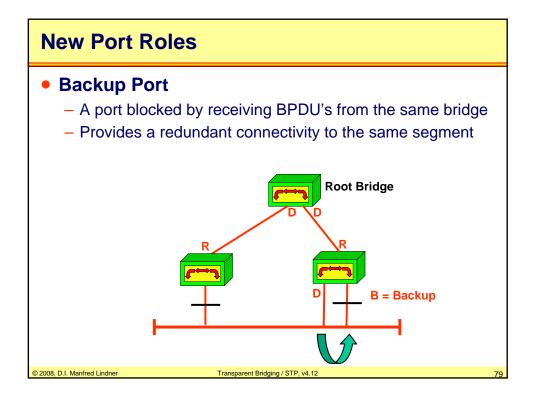
77

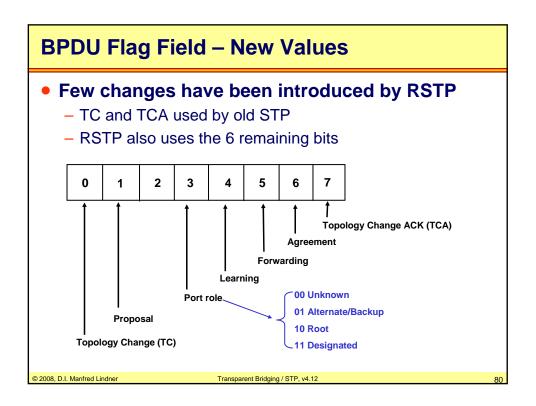
New Port Roles

Alternate Port Roles

- A port blocked by receiving BPDU's from a different bridge
- Provides an alternate path to the root bridge







NEW BPDU Handling

Faster Failure Detection

- BPDU's acting now as keepalives messages
 - Different to the 802.1D STP a bridge now sends a BPDU with its current information every <hello-time> seconds (2 by default), even if it does not receive any from the root bridge
- If hellos are not received for 3 consecutive times, port information is invalidated
 - because BPDU's are now used as keep-alive mechanism between bridges
- If a bridge fails to receive BPDU's from a neighbor, the connection has been lost
- No more max age and message age fields
 - · Hop count is used instead

© 2008, D.I. Manfred Lindner

Transparent Bridging / STP, v4.12

0.

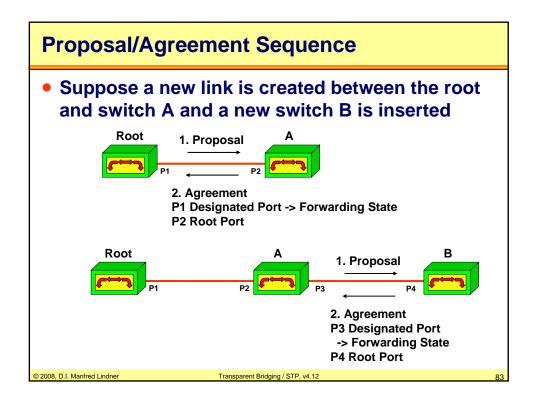
Proposal / Agreement

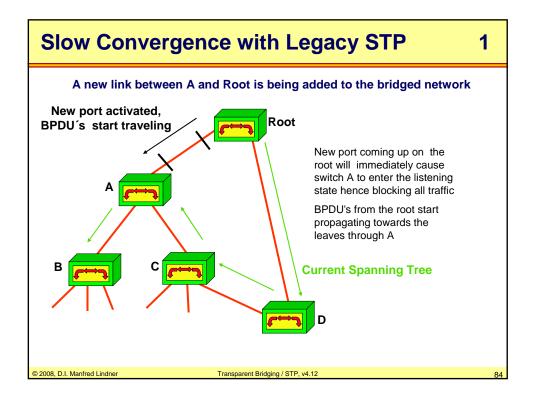
Explicit handshake between bridges

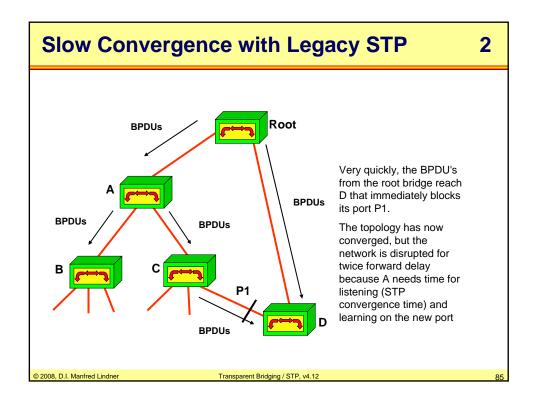
- Upon link up event the bridge sends a proposal to become designated for that segment
- Remote bridge responses with an agreement if the port on which it received the proposal is the root port of the remote bridge
- As soon as receiving an agreement, bridge moves the port to the forwarding state

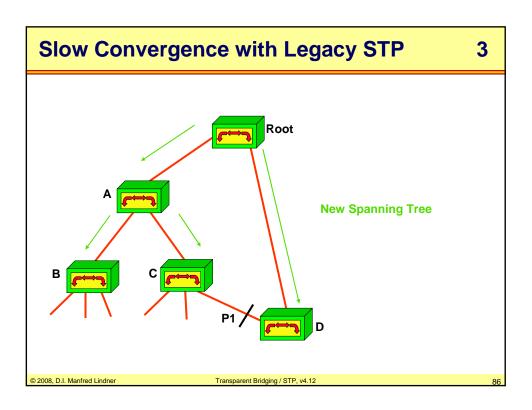
© 2008, D.I. Manfred Lindner

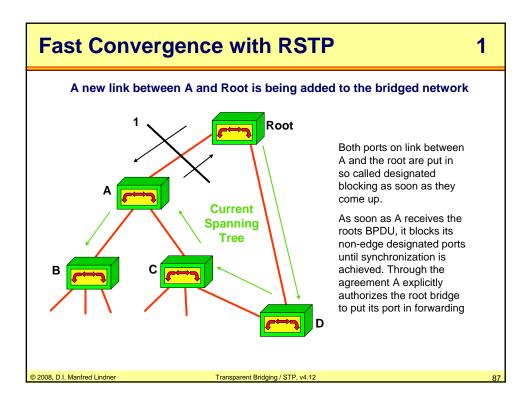
Transparent Bridging / STP, v4.12

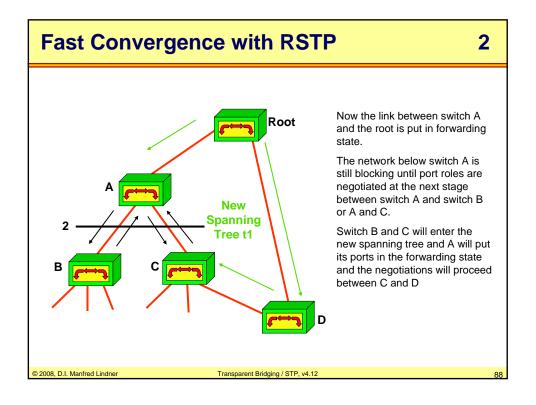


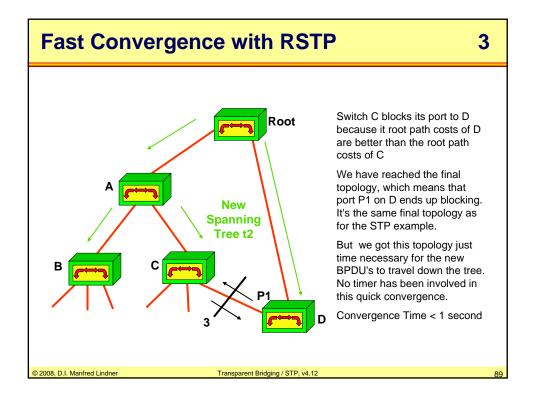












Rapid Transition to Forwarding State

- Most important feature in 802.1w
- The legacy STP was passively waiting for the network to converge before turning a port into the forwarding state
- New RSTP is able to actively confirm that a port can safely transition to forwarding
- Real feedback mechanism, that takes place between RSTP-compliant bridges
- To achieve fast convergence on a port, the protocol relies upon 2 new variables
 - Edge ports
 - Link type

© 2008, D.I. Manfred Lindner

Transparent Bridging / STP, v4.12

an

Rapid Transition to Forwarding State

Edge Ports

 Ports, which are directly connected to end stations cannot create bridging loops in the network and can thus directly perform on link setup transition to forwarding, skipping the listening and learning states

Link type

- Is automatically derived from the duplex mode of a port
 - A port operating in full-duplex will be assumed to be point-to-point
 - A port operating in half-duplex will be assumed to be a shared port

RSTP can only achieve rapid transition to forwarding

- On edge ports
- On point-to-point links (trunks between L2 switches)
- But not on shared Ports

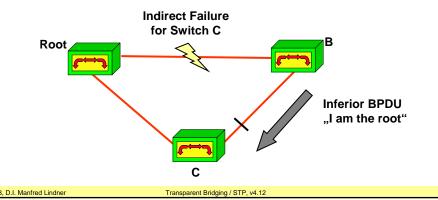
© 2008, D.I. Manfred Lindne

Transparent Bridging / STP, v4.12

01

Accepting Inferior BPDU's

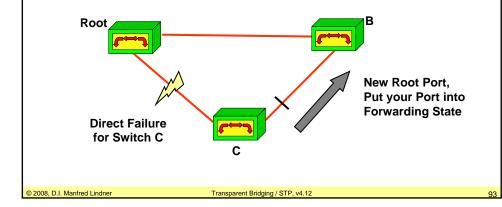
- B looses root port and sends BPDU claiming to be the root
- C immediately becomes designated for the blocked link between C and B and sends a proposal to B
- B sends an agreement and C set its port to forwarding
- Like Cisco's Backbone Fast



© 2008, D.I. Manfred Lindner

Accepting New Root Port BPDU's

- C looses root port and sends BPDU on the blocked link agreeing that this port is now root ort
- C set its port to forwarding
- Like Cisco's Uplink Fast



Topology Change Detection

- When a RSTP bridge detects a topology change, the following happens:
 - It starts the TC While timer with a value twice the hello time for all its non-edge designated ports and its root port if necessary
 - It flushes the MAC addresses associated with all these ports from the MAC bridging table

© 2008, D.I. Manfred Lindner

Transparent Bridging / STP, v4.12

Topology Change Propagation

- When a bridge receives a BPDU with the TC bit set from a neighbor, the following happens:
 - It clears the MAC addresses learnt on all its ports except the one that received the topology change
 - It starts the TC While timer and sends BPDU's with TC set on all its designated ports and root port (RSTP no longer uses the specific TCN BPDU, unless a legacy bridge need to be notified)

© 2008, D.I. Manfred Lindner

Fransparent Bridging / STP, v4.12

95

Topology Change Propagation The originator of the TC directly floods this information through the network Transparent Bridging / STP, v4.12 96

Agenda

- Introduction
- Transparent Bridging Basics
 - Principles
 - Broadcast Storm Problem
- Spanning Tree Protocol
 - Introduction
 - Details
 - Convergence
- Rapid Spanning Tree Protocol
- Routing Principles

© 2008, D.I. Manfred Lindner

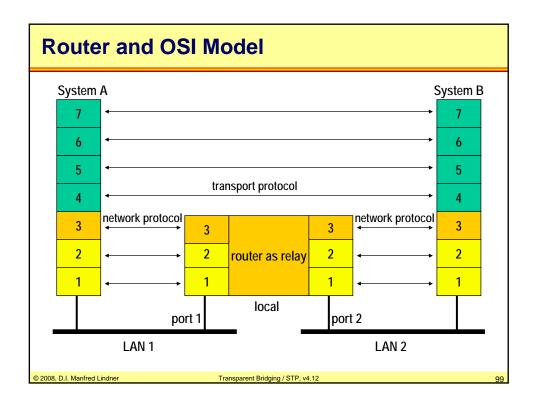
Transparent Bridging / STP, v4.12

Router

- router forwards packets
 - based on <u>layer 3 addresses</u> and protocols
- layer 3 address
 - structured versus unstructured layer 2 address
 - at least two level hierarchy: subnet and end system (host)
 - hardware independent
 - identifies a certain end system located in one subnet in a non-ambiguous way
 - a structured address is laid upon the unstructured MACaddress
- router connects
 - subnets knowing the best path to other subnets

© 2008, D.I. Manfred Lindner

Transparent Bridging / STP, v4.12



Requirements for Routers

consistent layer-3 functionality

- for entire transport system
- from one end-system over all routers in between to the other end-system
- hence routing is not protocol-transparent
 - all elements must speak the same "language"

end-system

- must know about default router
- on location change, end-system must adjust its layer 3 address

to keep the routing tables consistent

 routers must exchange information about the network topology by using routing-protocols

© 2008, D.I. Manfred Lindner

Transparent Bridging / STP, v4.12

Routing Facts

1

- in contrast to bridges
 - router maintains only the subnet-part of the layer 3 addresses in its routing table
 - the routing table size is direct <u>proportional to the number</u> of subnets and not to the number of end-systems
- transport on a given subnet
 - still relies on layer 2 addresses
- end systems forward data packets for remote destinations
 - to a selected router using the router's MAC-address as destination
 - only these packets must be processed by the router

© 2008, D.I. Manfred Lindner

Transparent Bridging / STP, v4.12

404

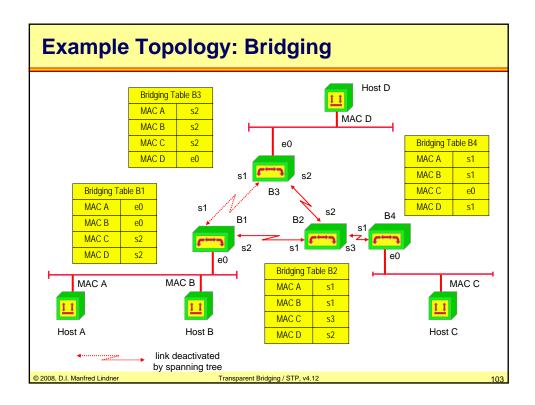
Routing Facts

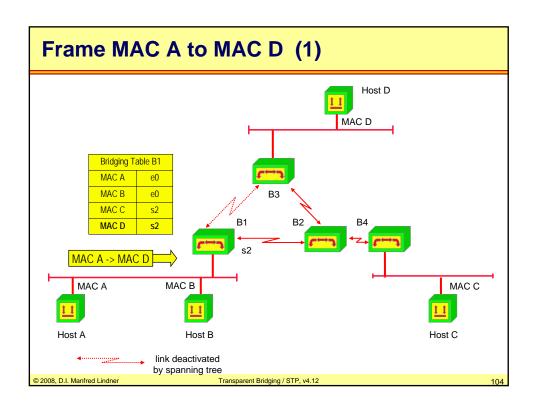
2

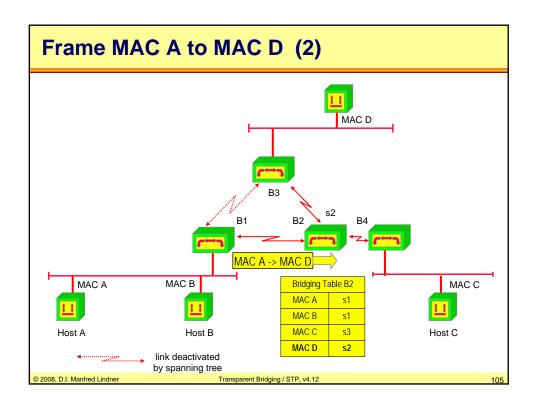
- flow control between router and end system is principally possible
 - end systems knows about the local router
- broadcast/multicast-packets in the particular subnet
 - are blocked by the router so broad/multicast traffic on the subnets doesn't stress WAN connections
- independent of layer 1, 2
 - so coupling of heterogeneous networks is possible
- routers can use redundant paths
 - meshed topologies are usual
- routers can use parallel paths for load balancing

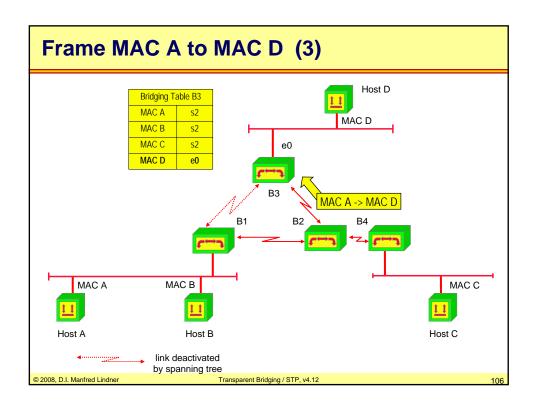
© 2008, D.I. Manfred Lindner

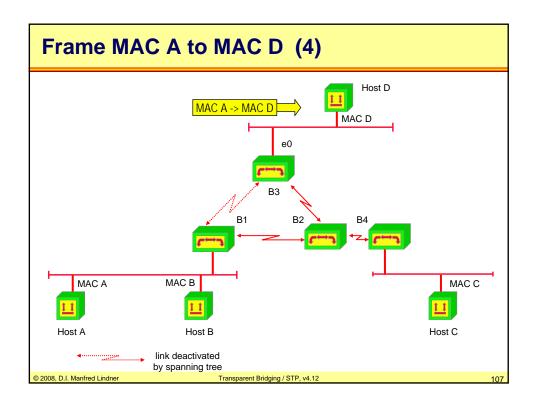
Transparent Bridging / STP, v4.12

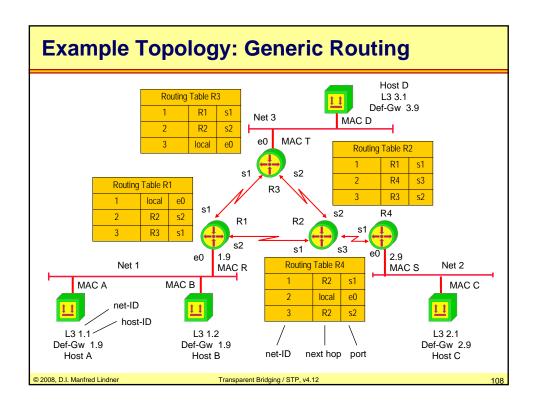


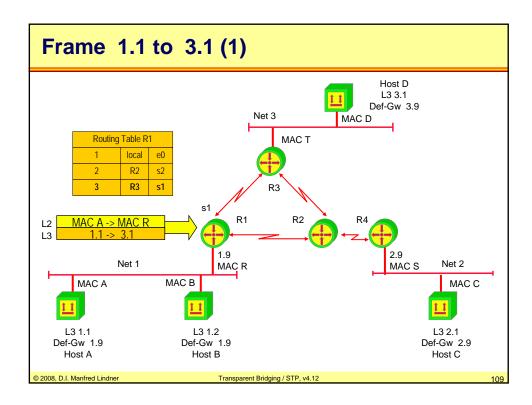


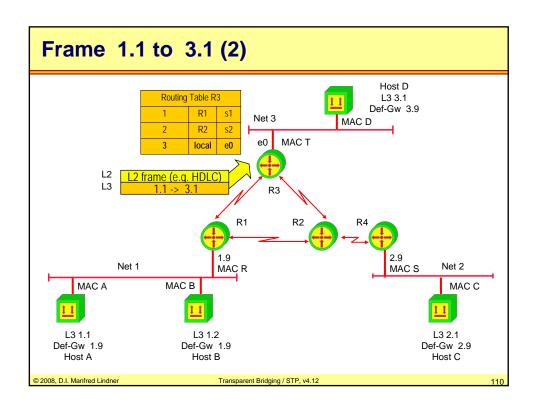


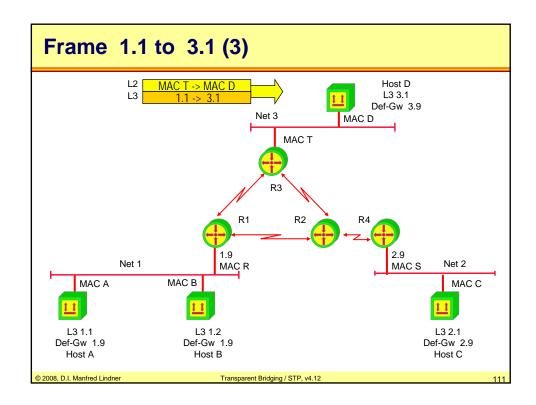


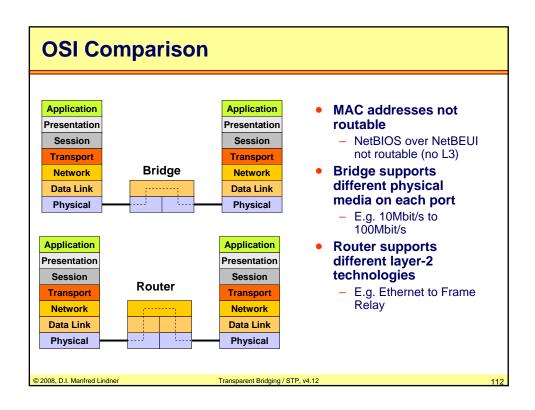












Transparent Bridging versus Routing Bridging Routing depends on MAC addresses requires structured addresses (must be configured) only invisible for end-systems; end system must know its transparent for higher layers default-router processes only frames must process every frame addressed to it number of table-entries = number of table-entries = number of all devices in the number of subnets only whole network Spanning Tree eliminates redundant lines and load redundant lines; no load balance balance possible no flow control (may be flow control is possible (router changed by usage of MAC is seen by end systems) Pause command)

Transparent Bridging versus Routing Bridging Routing no LAN/WAN coupling because does not stress WAN with of high traffic (broadcast subnet's broad- or multicasts; domain!) commonly used as "gateway" paths selected by STP may not router knows best way for each match communication behaviour frame / needs of end systems slower, because usually faster, because implemented in implemented in SW; address HW; no address resolution resolution (ARP) necessary location change of an endlocation change of an end-system does not require system requires adjustment of updating any addresses layer 3 address (may be solved by DHCP) Transparent Bridging / STP, v4.12

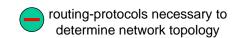
Transparent Bridging versus Routing

Bridging

Routing



spanning tree necessary against endless circling of frames and broadcast storms



- but today's L2 switches can improve network performance by enabling full duplex, collisionfree subnets
- MAC address resolution of routers is considered as unnecessary overhead -> trend goes to layer 3 switching (fast routing)

© 2008, D.I. Manfred Lindner

Transparent Bridging / STP, v4.12

445

Most Important!

- bridge separates LAN into
 - multiple collision domains
- but a <u>bridged</u> network is still
 - one broadcast domain
 - broadcast frames are always flooded
- <u>router</u> separates the whole Inter-network
 - into multiple broadcast domains
 - broadcast frames arriving on a LAN are always stopped at the router
 - they can't leave their origin LAN

© 2008, D.I. Manfred Lindner

Transparent Bridging / STP, v4.12

Ethernet Switching

- Ethernet switch is basically a bridge, differences are only:
 - faster because implemented in HW
 - multiple ports
 - improved functionality (e.g. VLAN)
 - note: TB is packet switching / connectionless mode with frame forwarding based on unstructured MAC addresses
- don't confuse it with WAN Switching!
 - e.g. X.25 switch, Frame Relay switch, ATM switch
 - completely different!
 - that is packet switching / connection-oriented mode with forwarding based on local connection identifiers

2008, D.I. Manfred Lindner

Transparent Bridging / STP, v4.12