



IA-32 Assembly

(Deeper and deeper...)

Emmanuel Fleury
LaBRI, Office 261
<emmanuel.fleury@labri.fr>





Outline

- Motivations & Warnings
 - IA-32 Processors Family
 - Memory Discipline: Stack & Heap
 - Computational Model
 - IA-32 Registers
 - IA-32 Assembly Syntax(es)
 - Basic Instruction Sets
 - Interruptions & System Calls
 - Advanced Instruction Sets
-



Motivations & Warnings



What's Assembly Useful For ?

- **Understand the machine**
(find bugs is easier, less program design errors, ...)
- **Code hardware dedicated subroutines**
(compilers, drivers, operating systems, ...)
- **Optimize high-level language code**
(manage your compiler...)
- **Reverse engineering**
(malware analysis, driver analysis, ...)

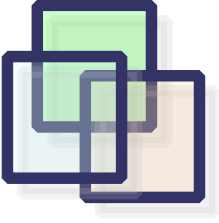
Knowing Assembly will enhance your code !



What's Assembly Bad At ?

- **Portability is lost**
(code is working only for only family of processors)
- **Optimization is tedious**
(compilers are more efficient than humans to optimize code)
- **Obfuscate the code**
(only few programmers can read assembly)
- **Debugging is difficult**
(most of the debuggers are lost when hitting assembly code)

Use it with caution and sparsity !!!



IA-32 Processors Family



Ancestors (Before IA-32)

- **Intel 4004 (1971): First microchip**
4bits memory words, 640b of addressable memory, 740kHz
- **Intel 8008 (1972):**
8bits memory words, 16kb of addressable memory, 800kHz
- **Intel 8086 (1978):**
16bits memory words, 1Mb of addressable memory, 10MHz
- **Intel 80286 (1982):**
16bits memory words, 16Mb of addressable memory, 12.5MHz





First IA-32 Generation

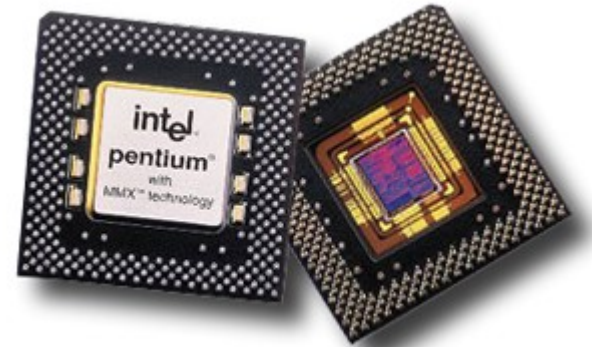
- **Intel 80386(DX) (1985):**
Introduce a Memory Management Unit (MMU)
32bits memory words, 4Gb of addressable memory, 16MHz
- **Intel 80486(DX) (1989):**
Mathematics co-processor built on-chip
32bits memory words, 4Gb of addressable memory, 16MHz





IA-32 Compatible Processors

- **Intel IA-32 Processors:**
Pentium, Pentium II, Pentium III, Pentium 4, Pentium M, Celeron, Core, Core 2;
- **AMD IA-32 Processors:**
K5, K6, K7 (Athlon, Sempron, Duron), K8 (Athlon, Sempron, Duron);
- **Transmeta IA-32 Processors:**
Crusoe, Efficeon.
- **Others...**

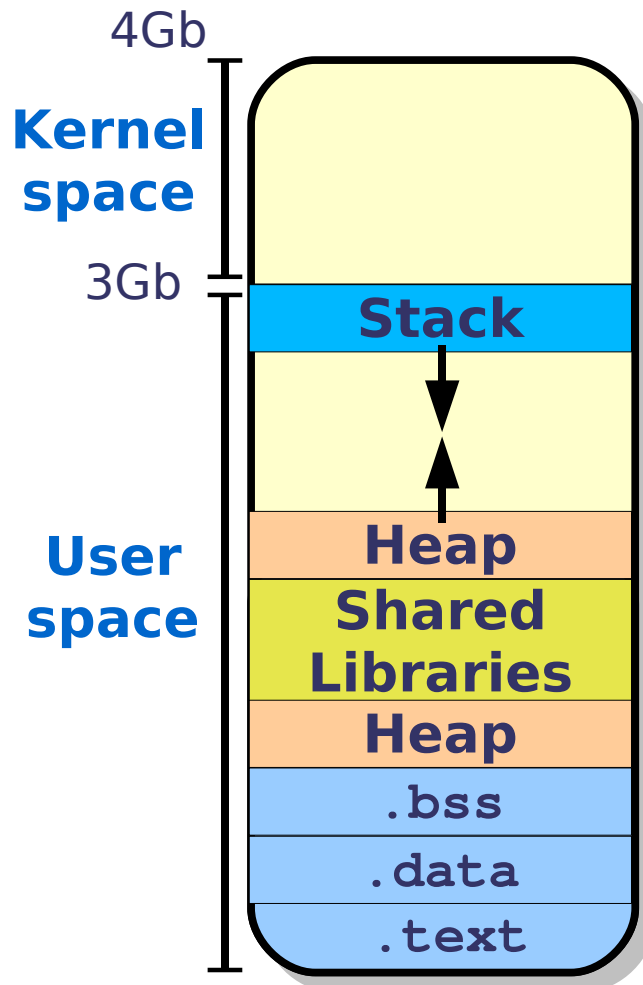




Memory Discipline: Stack & Heap



Linux Memory Layout

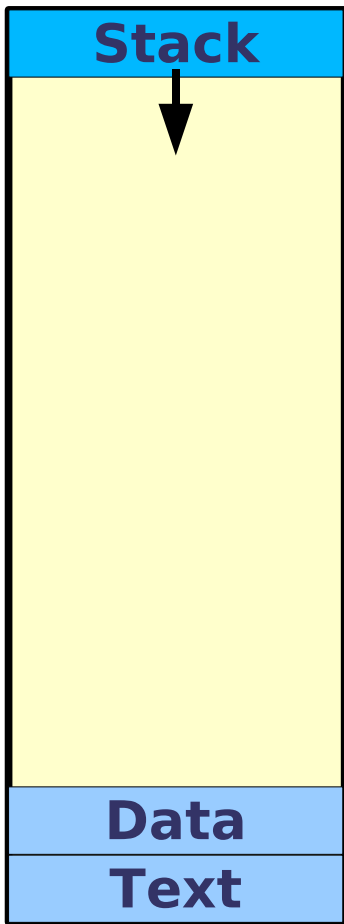


- **Stack**
Runtime stack (8Mb limit by default);
- **Heap**
Dynamically allocated storage (`malloc()`, `calloc()`, `new`, ...);
- **Shared/Dynamic Libraries**
Libraries routines (e.g. `printf()`, ...);
- **.bss, .data**
Statically allocated data (BSS start all zeroed);
- **.text, RODATA (Read-Only Data)**
Executable machines instructions and Read-Only DATA (string literals).

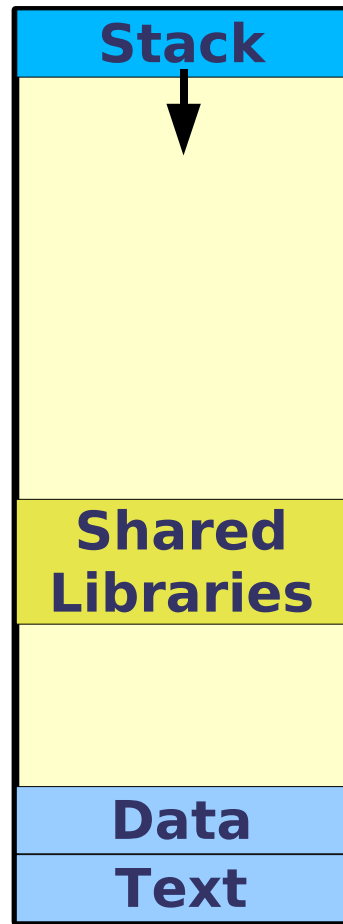


Linux Memory Allocation

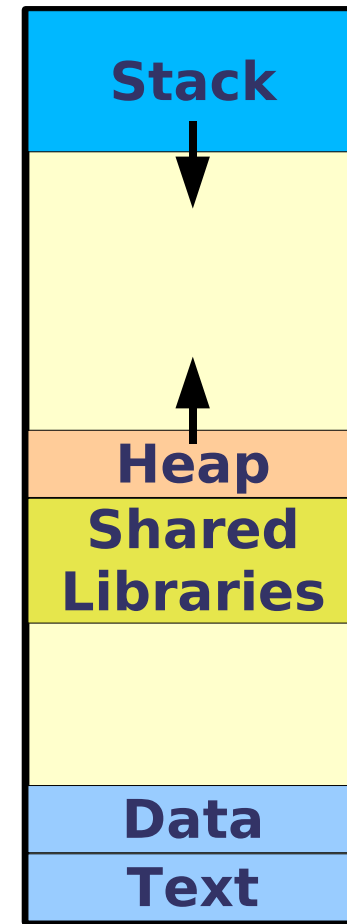
Initially



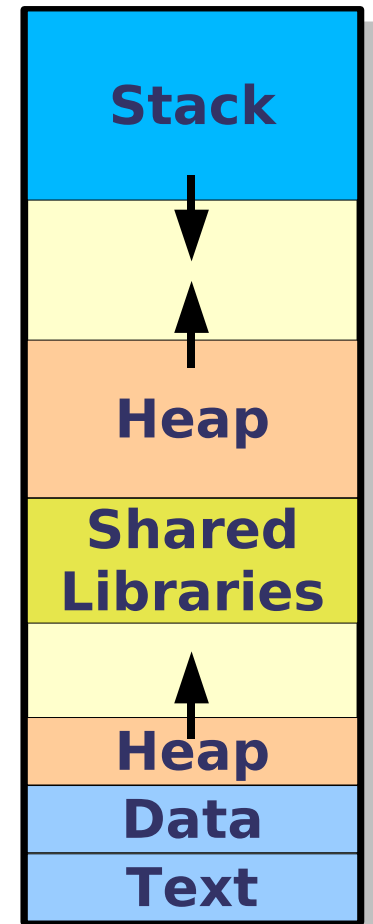
Linked



Few Heap



More Heap





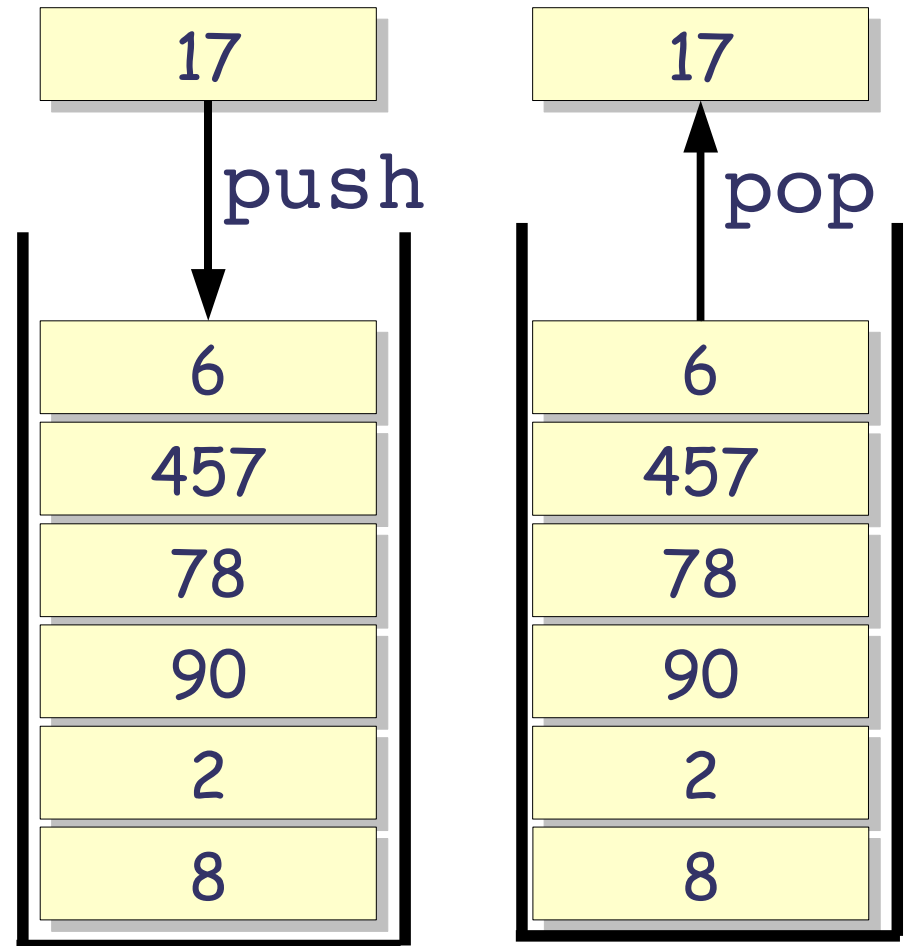
Stack Discipline



What is a Stack (FILO) ?

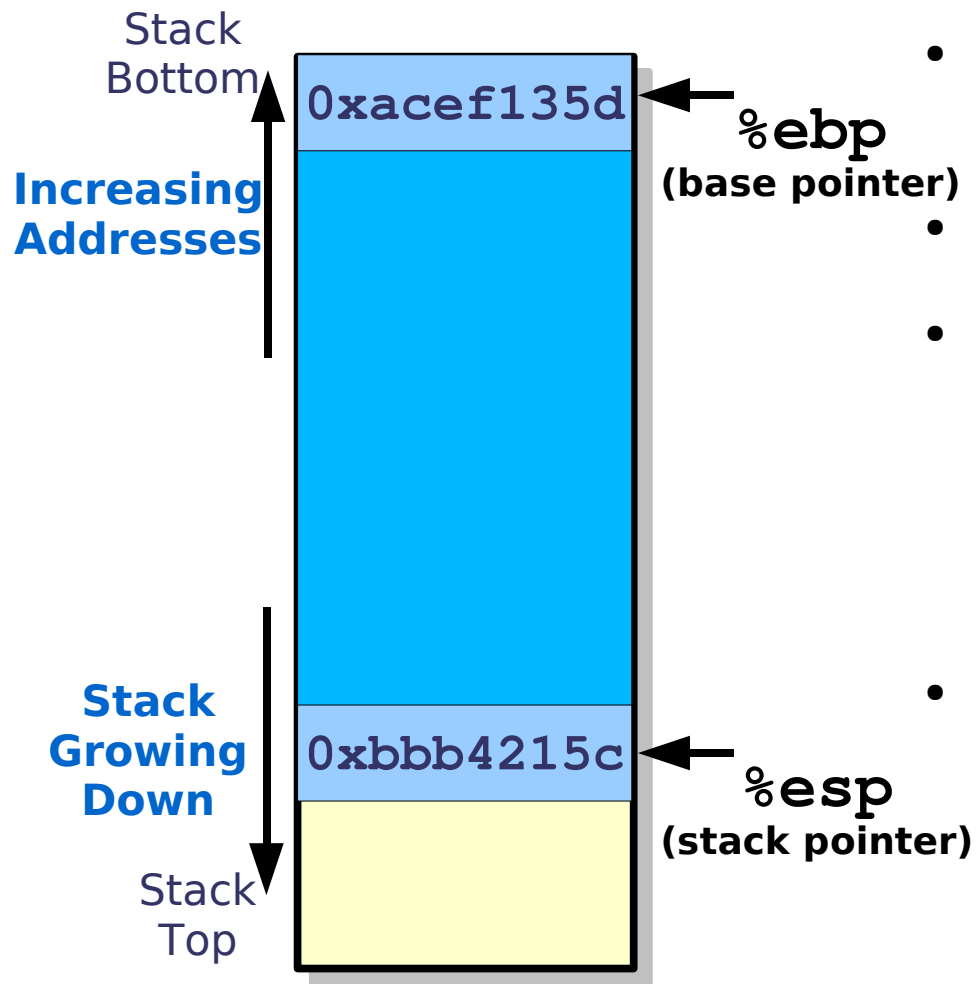
First In Last Out

- Implemented via:
Array, Linked-list
- Applications:
Stack based calculus
(CPU, parser, ...)
- Interface through:
Push & Pop





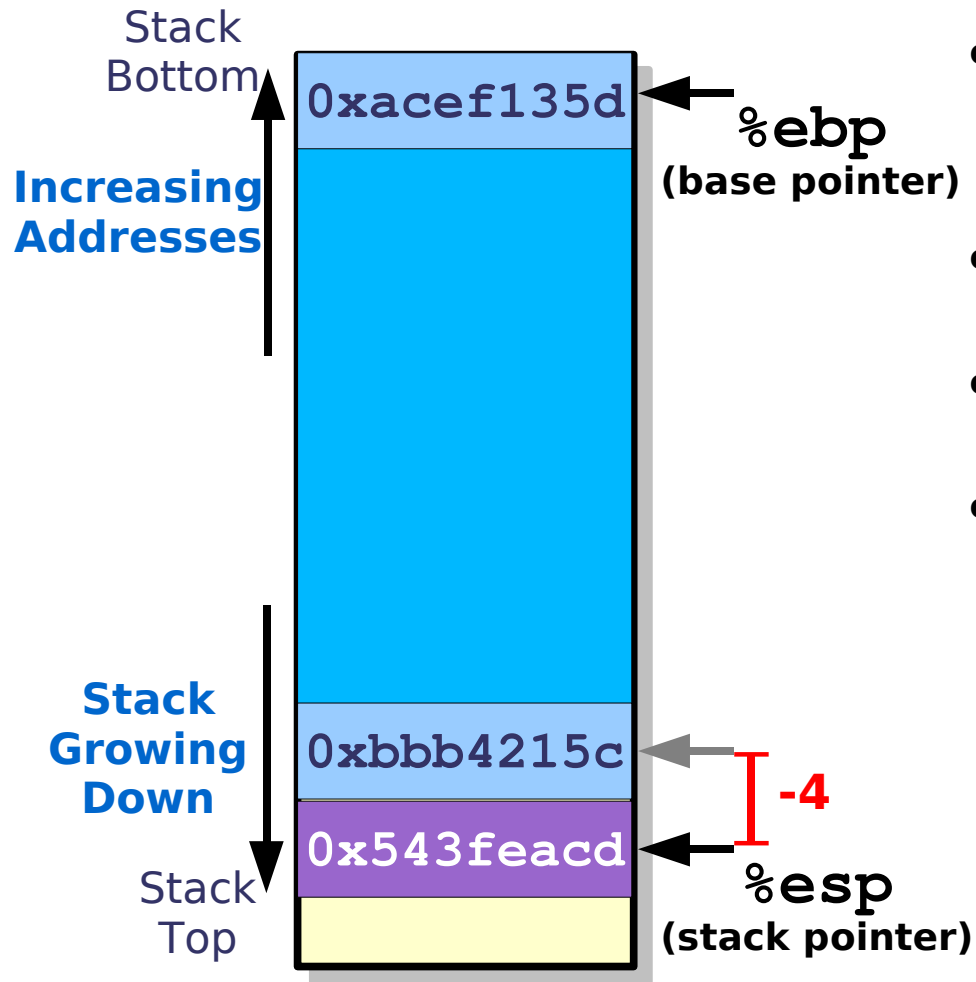
IA-32 Stack



- Memory zone managed with “*stack discipline*”
- Grows toward *lower addresses*
- Register `%esp` indicates lowest stack address
 - Mark the “top” element
 - “*stack pointer*” (SP)
- Register `%ebp` indicates highest stack address
 - Mark the “bottom” element
 - “*base pointer*” (BP)



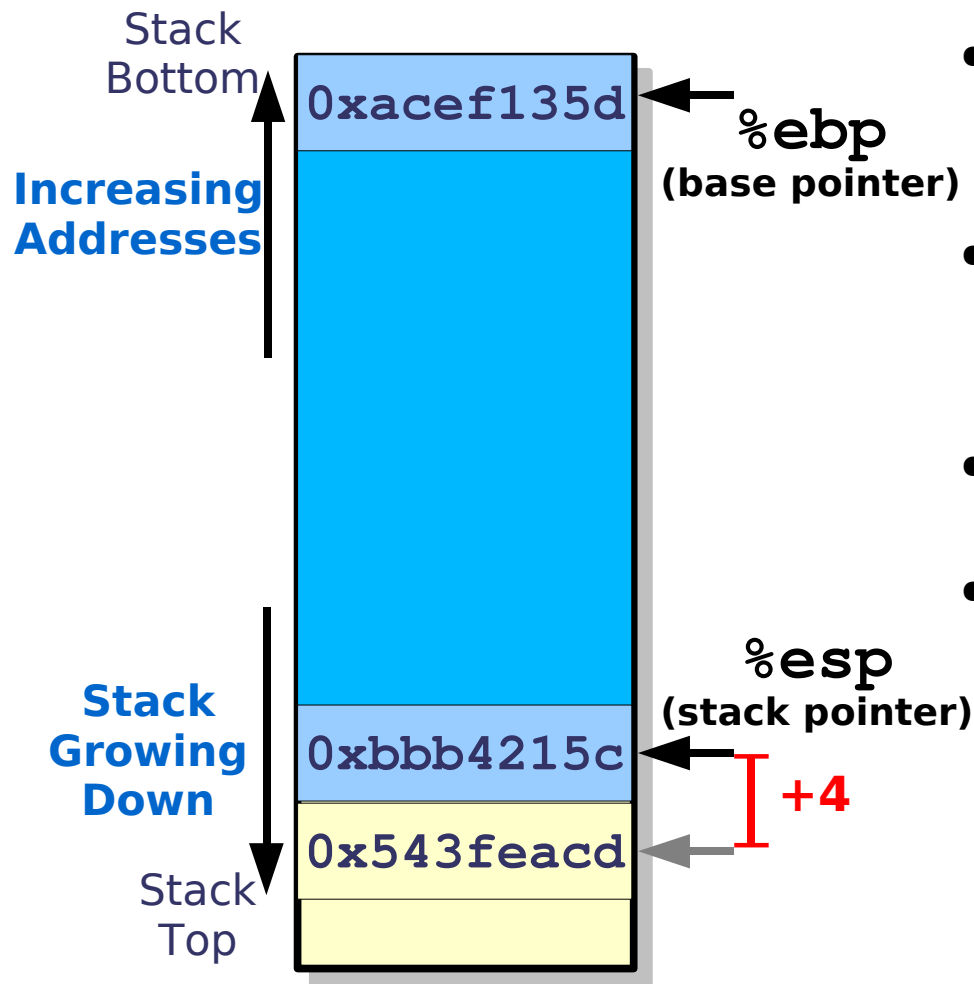
IA-32 Stack Pushing



- `pushl Src`
(Src=memory address)
- Fetch operand at Src
- Decrement `%esp` by 4
- Write operand at
(`%esp`)



IA-32 Stack Popping

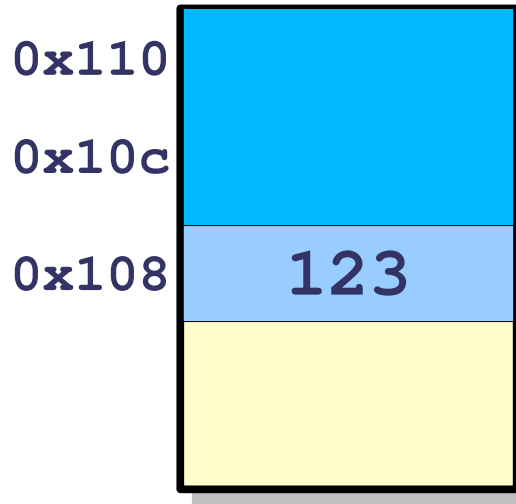


- `pop Dest`
(Dest=memory address)
- Read operand at
(`%esp`)
- Increment `%esp` by 4
- Write operand to
(Dest)



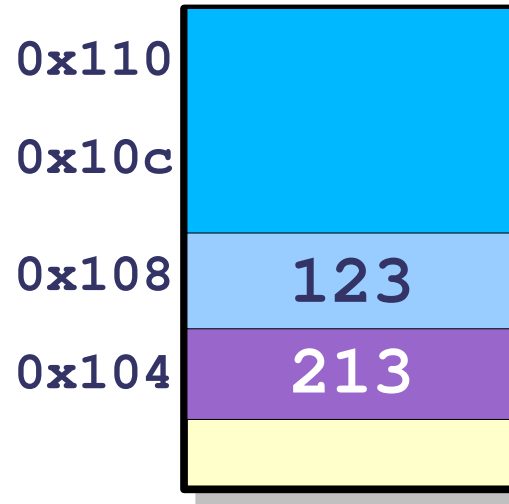
Stack Operation Examples

pushl %eax



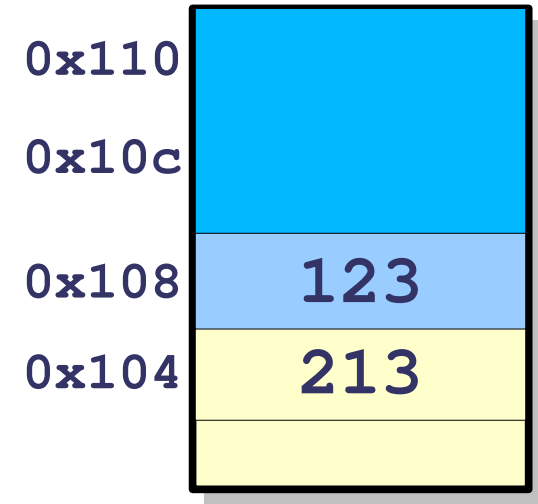
%eax	213
%edx	555
%esp	0x108
%ebp	0x110

pushl %eax

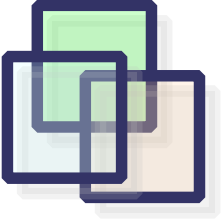


%eax	213
%edx	555
%esp	0x104
%ebp	0x110

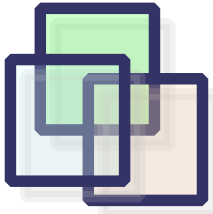
popl %edx



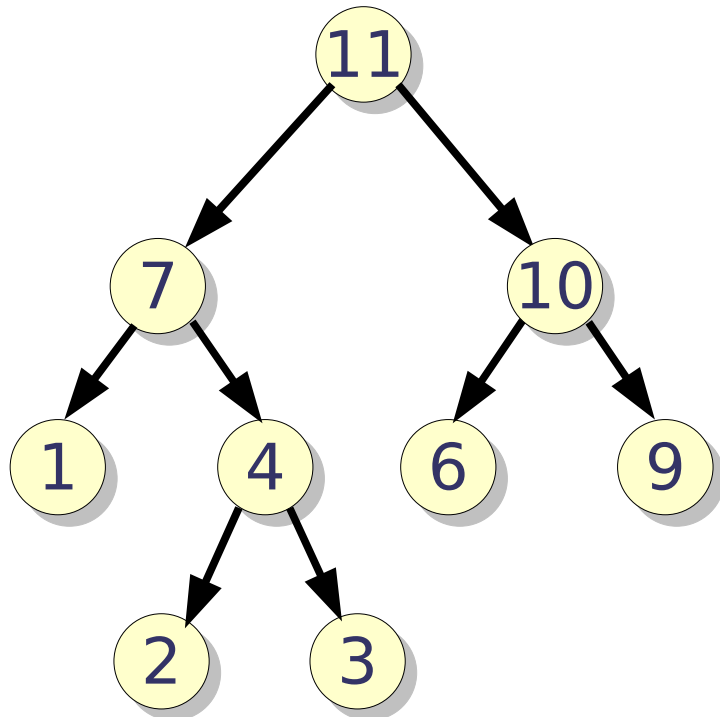
%eax	213
%edx	213
%esp	0x108
%ebp	0x110



Heap Discipline



What is a (binary) Heap ?



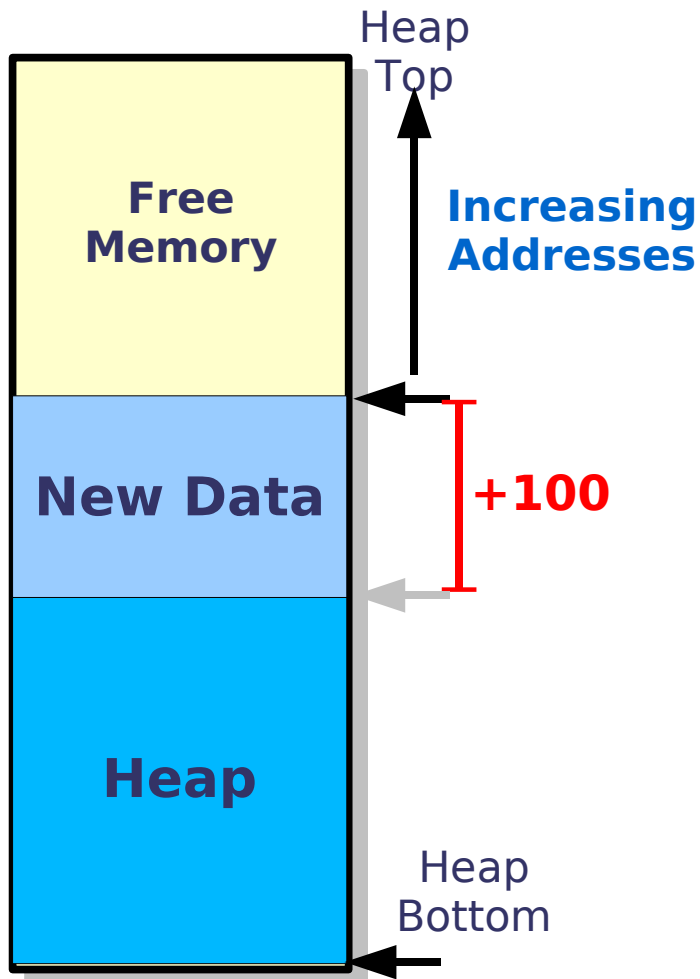
A **heap** is a tree structure such that, if A and B are nodes of a heap and B is a child of A, then:

$$\text{key}(A) \geq \text{key}(B)$$

- Implemented via:
 - Arrays
($a[i]$ has two children $a[2i+1], a[2i+2]$)
 - Trees
- Applications:
Quick access to data (databases)
- Groups of Data:
In Doug Lea malloc (dlmalloc) memory chunks are classified by size (bytes).



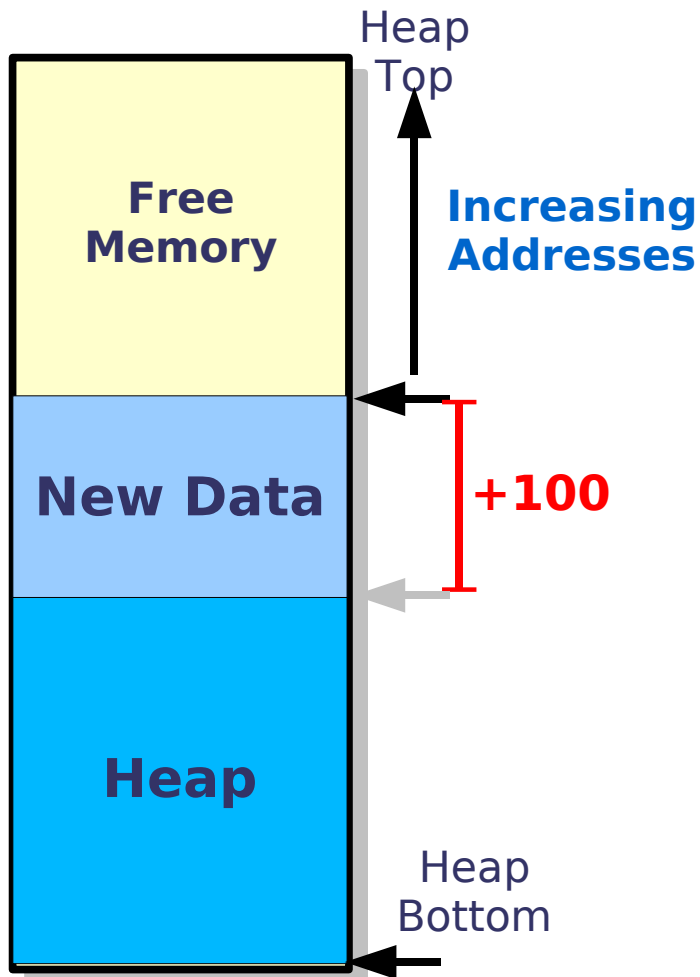
IA-32 Heap



- Memory zone managed with “*heap discipline*”
- Grows toward higher addresses
- From programmer point of view:
Managed through a language dependent interface (C, C++,...).
- From the system point of view:
Managed through specific system calls (`mmap()`, `brk()`).



IA-32 Heap (System Calls)



brk(void *end_data_segment) :

Sets the end of the data segment to the value specified by `end_data_segment`, when that value is reasonable, the system does have enough memory and the process does not exceed its max data size.

mmap() :

Asks to map length bytes starting at offset offset from the file (or other object) specified by the file descriptor `fd` into memory, preferably at address start. This latter address is a hint only, and is usually specified as 0. The actual place where the object is mapped is returned by `mmap()`, and is never 0.

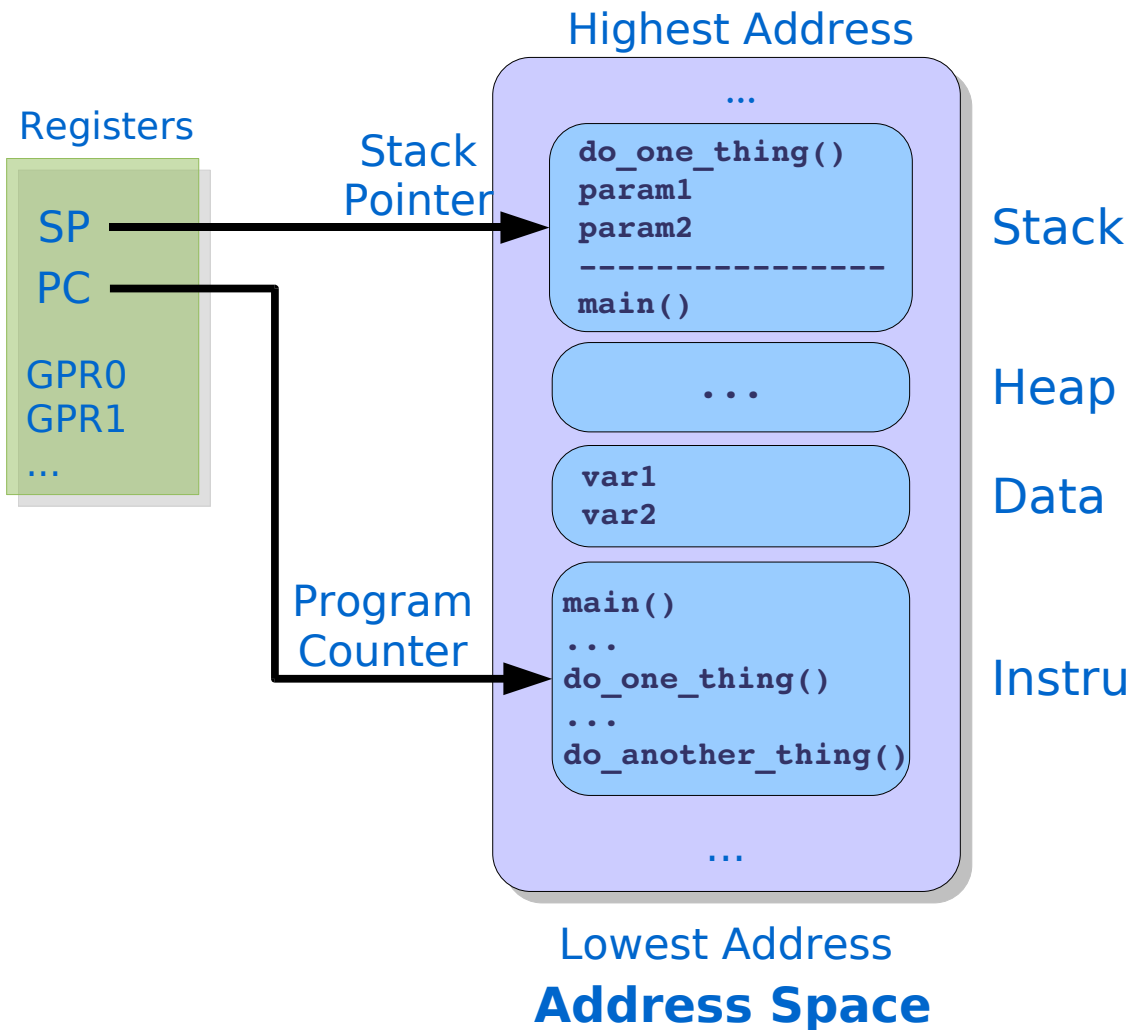
brk() is used for small chunks, where **mmap()** is used for big ones.



Computational Model



Computational Model



Registers:

- SP (Stack Pointer)
- PC (Program Counter)
- GPR (General Purpose Register)

Memory:

- Stack
- Heap
- Data
- Instructions



Computational Model (Example)

Computing “(5*3+2)-5”:

```
movl    $5, %eax    # eax = 5;  
imull   $3, %eax    # eax = 3*eax;  
addl    $2, %eax    # eax = eax+2;  
subl    $5, %eax    # eax = eax-5;
```

Instructions

“Immediate”
Values

Register
Name

Comments



Computational Model (Example)

Computing “(5*3+2)-5”:

```
→ movl    $5, %eax    # eax = 5;  
   imull   $3, %eax    # eax = 3*eax;  
   addl    $2, %eax    # eax = eax+2;  
   subl    $5, %eax    # eax = eax-5;
```

eax

5



Computational Model (Example)

Computing “(5*3+2)-5”:

```
movl    $5, %eax    # eax = 5;  
→ imull  $3, %eax    # eax = 3*eax;  
addl    $2, %eax    # eax = eax+2;  
subl    $5, %eax    # eax = eax-5;
```

eax 5 *3=15



Computational Model (Example)

Computing “(5*3+2)-5”:

```
movl    $5, %eax    # eax = 5;
imull   $3, %eax    # eax = 3*eax;
→ addl   $2, %eax    # eax = eax+2;
subl    $5, %eax    # eax = eax-5;
```

eax **15** +2=17



Computational Model (Example)

Computing “(5*3+2)-5”:

```
movl    $5, %eax    # eax = 5;
imull    $3, %eax    # eax = 3*eax;
addl    $2, %eax    # eax = eax+2;
→ subl    $5, %eax    # eax = eax-5;
```

eax **17** -5=12



Computational Model (Example)

Computing “(5*3+2)-5”:

```
movl    $5, %eax    # eax = 5;
imull    $3, %eax    # eax = 3*eax;
addl    $2, %eax    # eax = eax+2;
subl    $5, %eax    # eax = eax-5;
```

eax 12



IA-32 Registers

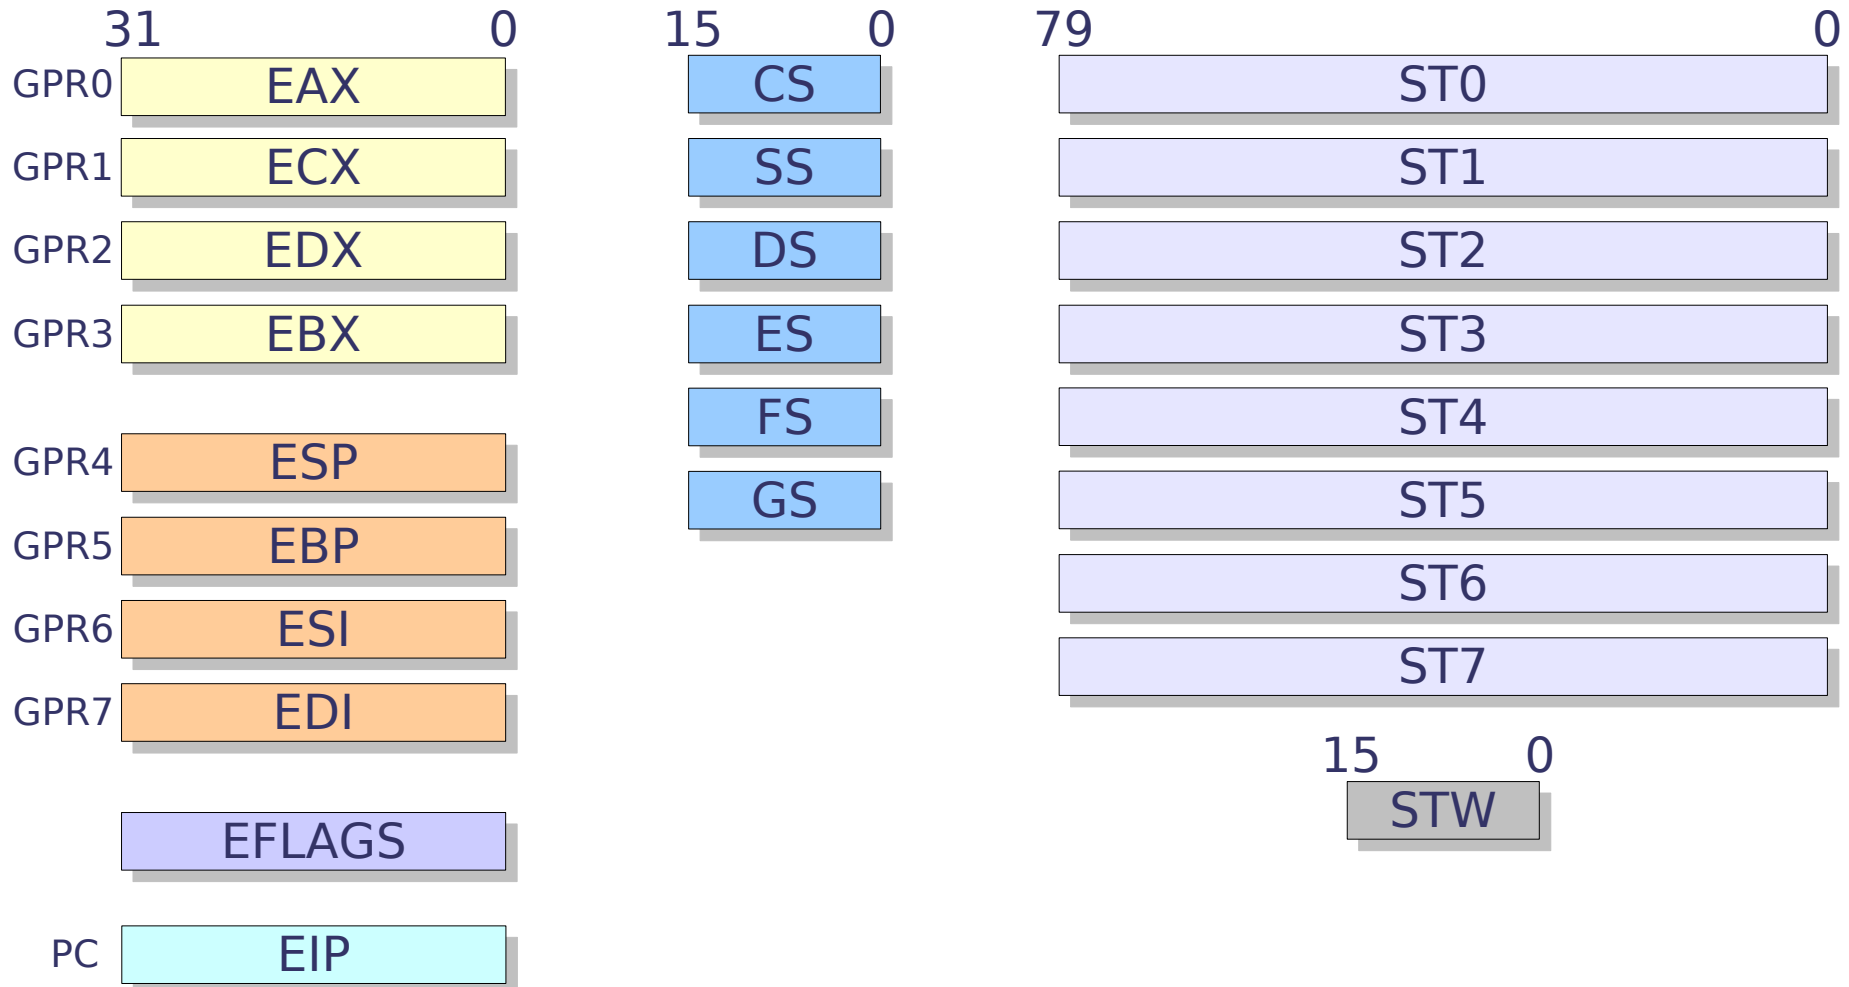


IA-32 Registers

- **Data Registers (Read/Write)**
(EAX, EBX, ECX, EDX)
 - **Index & Pointers Registers (Read/Write)**
(EBP, EIP, ESP, ESI, EDI)
 - **Segment Registers (Protected)**
(CS, DS, ES, FS, GS, SS)
 - **Flags Registers (Read)**
(EFLAGS)
 - **Floating-point Registers (Read/Write)**
(ST0, ..., ST7)
-

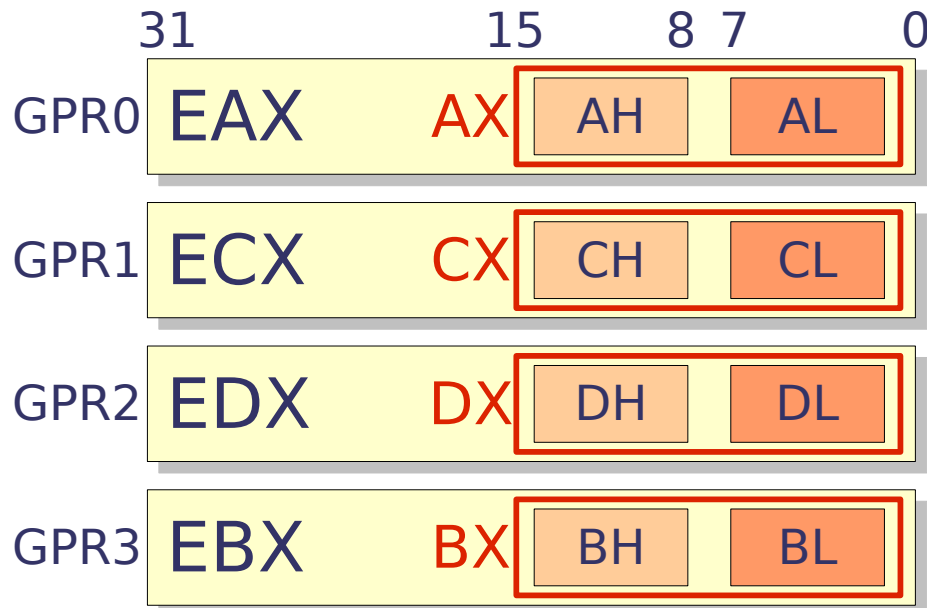


IA-32 Registers





Data Registers



For **backward compatibility** old **8 bits** and **16 bits** registers have been **preserved**. You still can address them.

EAX (Accumulator)

Accumulator for operands and results data (addition, subtraction, ...);

EBX (Base Register)

Usually used to store the base address of a data-structure in memory;

ECX (Count Register)

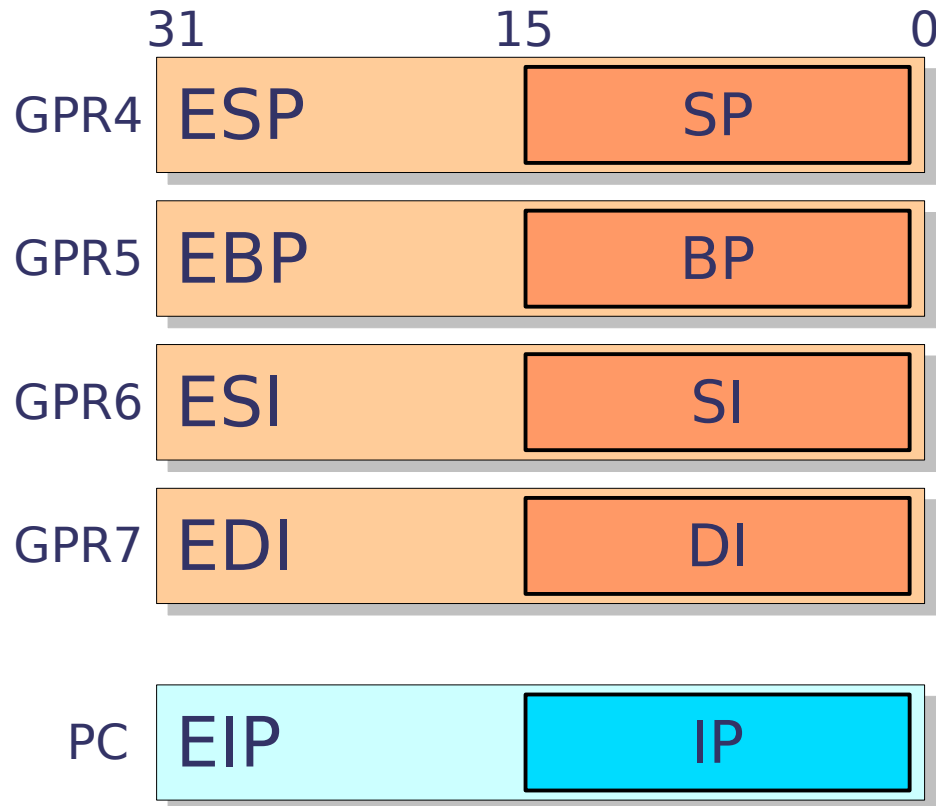
Usually used as a loop counter;

EDX (Data Register)

Used to store operand and result for multiplications and divisions.



Index & Pointers Registers



For **backward compatibility** old **16 bits** registers have been **preserved**.
You still can address them.

ESP (Stack Pointer)

Pointer to the last cell of the stack;

EBP (Base Pointer)

Pointer to the base cell of the current stack frame;

ESI (Source Index)

Used in string operations as source;

EDI (Destination Index)

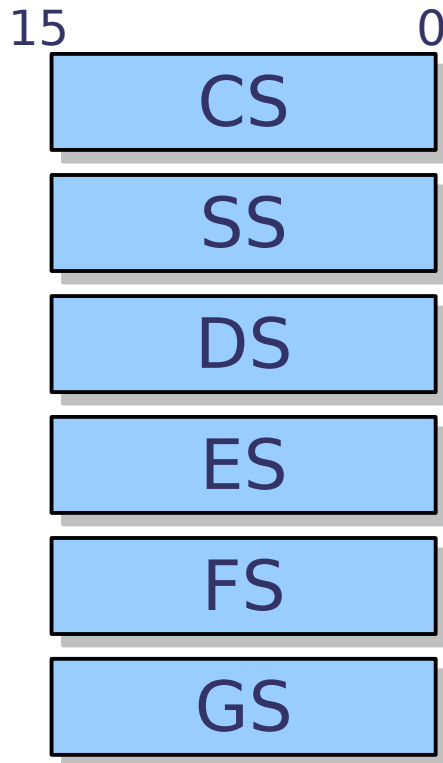
Used in string operations as destination;

EIP (Instruction Pointer)

Point to the next instruction.



Segment Registers



CS (Code Segment)

Point to the current code-segment;

SS (Stack Segment)

Point to the current stack-segment;

DS (Data Segment)

Point to the current data-segment;

ES, FS, GS (Extra Data Segments)

Extra segments registers available for far pointer addressing (video memory and others).



Flag Registers (Types)

EFLAGS

FLAGS

Status Flags (stat)

Give the result of arithmetic instructions (add, sub, mul or div)

Control Flags (ctrl)

Change the behavior of the processor on some instructions (std, cld)

System Flags (sys)

Accessible by kernel only



Flag Registers (FLAGS)



CF (Carry Flag, [stat](#)): Left most bit of result; **IF** (Interrupt Flag, [sys](#)): '1' if interruptions are on;

PF (Parity Flag, [stat](#)): '1' if result is even; **DF** (Direction Flag, [ctrl](#)): Strings reading order. Modified by [std](#) ('1', go from higher to lower addresses) and [cld](#) ('0', reverse order);

AF (Auxiliary Carry Flag, [stat](#)): Extra carry flag; **OF** (Overflow Flag, [stat](#)): '1' if an overflow occurs;

ZF (Zero Flag, [stat](#)): '1' if result is '==0';

SF (Sign Flag, [stat](#)): '1' if result is '<0';

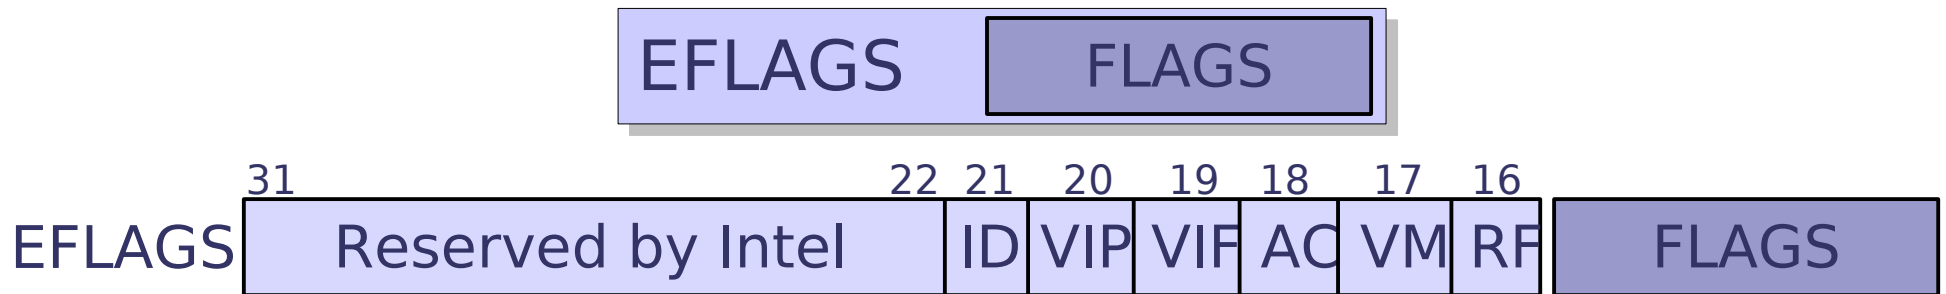
TF (Trap Flag, [sys](#)): Set CPU in single-step mode;

IOPL (I/O Privilege Level, [sys](#)): Current task privilege;

NT (Nested Task Flag, [sys](#)): '1' if current task is linked to the previous one.



Flag Registers (EFLAGS)



RF (Resume Flag, sys):
Set CPU in debug mode;

VM (Virtual Mode, sys):
Set the 8086 virtual mode (unset the protected mode);

AC (Alignment Check Flag, sys):
Set alignment checking mode for memory references;

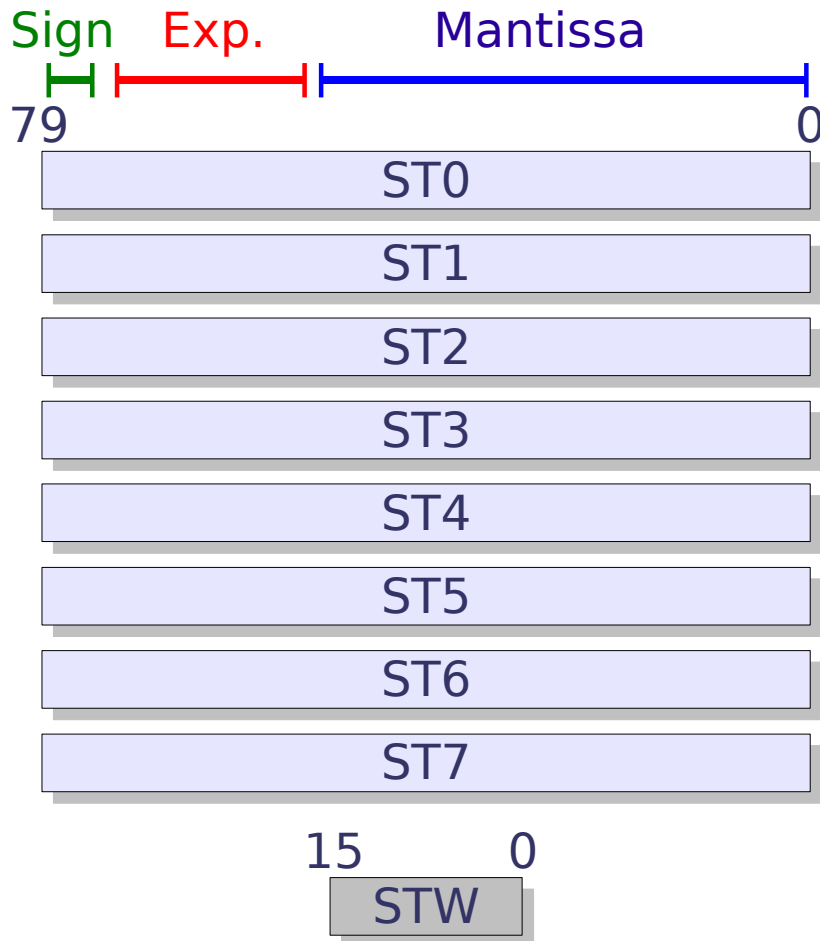
VIF (Virtual Interrupt Flag, sys):
Virtual image of the IF flag (used in conjunction with VIP);

VIP (Virtual Interrupt Pending Flag, sys):
Indicates pending interruptions ('1' if one is pending);

ID (ID Flag, sys):
Triggers the CPUID instruction support.



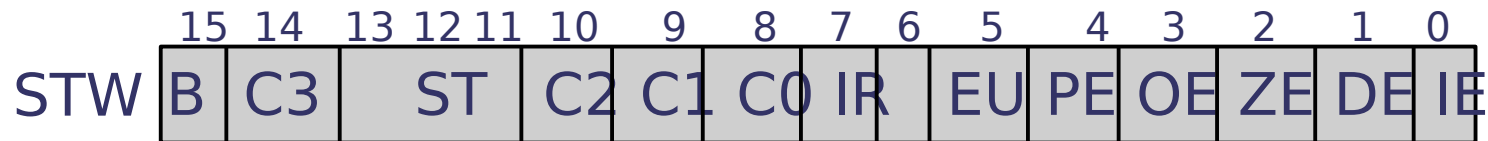
Floating-point Registers



- 80 bits wide registers
- Accessed as a stack (can't be addressed directly)
- Decomposed into:
 - Sign: 1 bit;
 - Exponent: 15 bits;
 - Mantissa: 64 bits;



Status Word Register



Exception Flags (bits 0-6)

- IE** Invalid Operation Exception
(Uncertain result)
- DE** Denormalized Operation Exception
(Operation on denormalized number)
- ZE** Zero Divide Exception (Division by zero)
- OE** Overflow Exception (Result is too large)
- UE** Underflow Exception (Result is too small)
- PE** Precision Exception
(Loss of precision occurred)

ST (Stack pointer): Counting free registers;

Condition Code (C0,C1,C2,C3):
Store results of boolean tests on float in these bits;

Others (IR, B):

- IR** Interruption Request
- B** Busy (on-going operation)



IA-32 Assembly Syntax(es)



Intel vs. AT&T Syntaxes

Tools are using two types of syntaxes !

- **Intel Syntax:**

All the Intel manuals are written in this syntax. A lot of tools are using this syntax, mostly tools from the *“Microsoft community”*.

- **AT&T Syntax:**

All the GNU binutils tools use this syntax. Most of the *“Unix community”* is also using it.



Intel vs. AT&T

(Direction of operand)

Intel Syntax:

The **first** operand is the **destination** and the **second** operand is the **source**.

Intel Syntax

Instr.	dest, src
mov	eax, [ecx]

AT&T Syntax:

The **first** operand is the **source** and the **second** operand is the **destination**.

AT&T Syntax

Instr.	src, dest
mov	(%ecx), %eax



Intel vs. AT&T

(Data types)

Intel Syntax:

- **10d**: Decimal value (**10** is ok);
- **10h**: Hexadecimal value;
- **1**: Immediate value;
- **eax**: Register;
- **byte ptr**: Address of a byte (8bits)
- **word ptr**: Address of a word (16bits)
- **dword ptr**: Address of a long (32bits)

AT&T Syntax:

- **0d10**: Decimal value (**10** is ok);
- **0x10**: Hexadecimal value;
- **\$1**: Immediate value;
- **%eax**: Register;
- **movb**: Operand on bytes (8bits)
- **movw**: Operand on words (16bits)
- **movl**: Operand on longs (32bits)



Intel vs. AT&T

(Data types)

Intel Syntax

```
mov    eax, 1
mov    ebx, 0ffh
int    80h
mov    al, bl
mov    ax, bx
mov    eax, ebx
mov    eax, dword ptr [ebx]
```

AT&T Syntax

```
movl    $1, %eax
movl    $0xff, %ebx
int      $0x80
movb    %bl, %al
movw    %bx, %ax
movl    %ebx, %eax
movl    (%ebx), %eax
```



Intel vs. AT&T

(Addressing memory)

Intel Syntax:

The address is enclosed in brackets ("**[**", "**]**") and specified by:

[base+index*scale+offset]

Intel Syntax

```
mov  eax, [ebx]
mov  eax, [ebx+3]
mov  eax, [ebx+20h]
add  eax, [ebx+ecx*2h]
lea  eax, [ebx+ecx]
sub  eax, [ebx+ecx*4h-20h]
```

AT&T Syntax:

The address is enclosed in parenthesis ("**(**", "**)**") and specified by:

offset (base, index, scale)

AT&T Syntax

```
movl  (%ebx), %eax
movl  3(%ebx), %eax
movl  0x20(%ebx), %eax
addl  (%ebx, %ecx, 0x2), %eax
leal  (%ebx, %ecx), %eax
subl  -0x20(%ebx, %ecx, 0x4), %eax
```



Basic Instruction Sets



Types of Instructions

Memory Operators

Moving data, stack management;

Arithmetic & Logic

Bitwise operators, shift & rotate, Integer arithmetic, floating-point arithmetic;

Flow Control

Tests, jump, loops;



Moving Data

Mnemonic	Operand	Operand	Operation
mov	src	dst	src → dst
xchg	dst	dst	dst ↔ dst
lea	mem	reg	mem → reg
movz	src8	reg16	zero-extended src → reg
	src8	reg32	
	src16	reg32	
movsz	src8	reg16	sign-extended src → reg
	src8	reg32	
	src16	reg32	
cbw	-	-	sign-extended %al → %ax
cwd	-	-	sign-extended %ax → %edx.%ax
cdq	-	-	sign-extended %eax → %edx.%eax
cwde	-	-	sign-extended %ax → %eax

lea = Load Effective Address



Moving Data (Example)

- **The code:**

```
# Example 1
```

```
.globl main
```

```
main:
```

```
    movl    $20, %eax
```

```
    ret
```

- **Building the binary:**

```
gcc -o example1 asm.s
```

- **Running the binary:**

```
./example1
```



Stack Management

Mnemonic	Operand	Operation
push	src8	$\%esp-1 \rightarrow \%esp$; $src \rightarrow (\%esp)$
	src16	$\%esp-2 \rightarrow \%esp$; $src \rightarrow (\%esp)$
	src32	$\%esp-4 \rightarrow \%esp$; $src \rightarrow (\%esp)$
pop	dst8	$(\%esp) \rightarrow dst$; $\%esp+1 \rightarrow \%esp$
	dst16	$(\%esp) \rightarrow dst$; $\%esp+2 \rightarrow \%esp$
	dst32	$(\%esp) \rightarrow dst$; $\%esp+4 \rightarrow \%esp$
pushf	-	$\%esp-4 \rightarrow \%esp$; $\%eflags \rightarrow (\%esp)$
popf	-	$(\%esp) \rightarrow \%eflags$; $\%esp+4 \rightarrow \%esp$
pusha	-	Push $\%eax, \%ecx, \%edx, \%ebx, \%esp, \%ebp, \%esi, \%edi$
popa	-	Pop $\%eax, \%ecx, \%edx, \%ebx, \%esp, \%ebp, \%esi, \%edi$
enter	-	Create a stack frame
leave	-	Restore the previous stack frame



Push'n Pop

```
.glob main
```

```
main:
```

```
    movl    $20, %eax
```

```
    pushl   %eax
```

```
    popl    %ebx
```

```
# Push in the stack
```

```
# Pop from the stack
```

```
    movl    $15, -4(%ebp)
```

```
    movl    4(%ebp), %ebx
```

```
# Push in the stack
```

```
# Pop from the stack
```

```
    ret
```



Memory Allocation

```
.glob main
```

```
main:
```

```
    # Allocate memory space
```

```
    pushl %ebp           # Save base pointer
```

```
    movl  %esp, %ebp     # Set stack pointer at base pointer
```

```
    subl  $8, %esp       # Allocate memory space for two words
```

```
    # Data manipulations
```

```
    pushl $10            # Push 10 in the stack
```

```
    pushl $15            # Push 15 in the stack
```

```
    popl  %eax           # Pop 15 from the stack
```

```
    popl  %ebx           # Pop 10 from the stack
```

```
    # Leaving stack-frame
```

```
    movl  %ebp, %esp     # Restore previous stack-pointer
```

```
    popl  %ebp           # Restore the old base pointer
```

```
    ret                 # Restore previous execution flow
```



Always Restore the Stack

```
.glob main
```

```
main:
```

```
    # Data manipulations
```

```
    pushl $10           # Push 10 in the stack
```

```
    pushl $15           # Push 15 in the stack
```

```
    popl  %eax          # Pop 15 from the stack
```

```
    popl  %ebx          # Pop 10 from the stack
```

```
    # Stack-pointer is not restored
```

```
    ret                # Restore previous execution flow
```

Running the program:

```
#bash> ./test  
Segmentation fault
```



Bitwise Operators

Mnemonic	Operand	Operand	Operation	Touched Flags
and	src	dst	src & dst → dst	SF,ZF,PF
or	src	dst	src dst → dst	
xor	src	dst	src ^ dst → dst	
test	src	dst	src & dst (result discarded)	
not	dst		~dst → dst	-
cmp	src	dst	sub src,dst (result discarded)	OF,SF,ZF,AF,CF,PF

Example: `.globl main`
`main:`
 `movl $8, %eax`
 `andl $9, %eax`
 `notl %eax`
L0: `cmp $8, %eax` `# %eax == 8`
 `jnz L1` `# Jump to L1 if ZF<>0`
 `jmp L0` `# Jump to L0`
L1: `ret`



Shift & Rotate

Mnemonic	Operand	Operand	Operation	Touched Flags
shl/sal	src	dst	left shift dst of src bits	CF,OF
shr/sar	src	dst	right shift dst of src bits	
rol	src	dst	left rotate dst of src bits	
ror	src	dst	right rotate dst of src bits	

Note: **shl/shr** are unsigned version of **sal/sar**.

2⁷ Multiplication:

```
.globl main
main:
    shl    $7, %eax    # %eax*2^7
    ret
```



Addition/Subtraction (Int)

Mnemonic	Operand	Operand	Operation	Touched Flags
add	src	dst	$\text{src} + \text{dst} \rightarrow \text{dst}$	OF, SF, ZF, AF, CF, PF
adc	src	dst	$\text{src} + \text{dst} + \text{CF} \rightarrow \text{dst}$	
sub	src	dst	$\text{dst} - \text{src} \rightarrow \text{dst}$	
sbb	src	dst	$\text{dst} - \text{src} - \text{CF} \rightarrow \text{dst}$	
inc	dst		$\text{dst} + 1 \rightarrow \text{dst}$	OF, SF, ZF, AF, PF
dec	dst		$\text{dst} - 1 \rightarrow \text{dst}$	
neg	dst		$-\text{dst} \rightarrow \text{dst}$	OF, SF, ZF, AF, PF CF=0 if $\text{dst}=0$

Example: `.globl main`
`main:`

```
    movl    $15, %eax
    subl    $7,  %eax    # %eax = %eax-7
    addl    $30, %eax    # %eax = %eax+30
    decl    %eax         # %eax = %eax-1
    ret
```



Multiplication/Division (Int)

Mnemonic	Operand	Operation	Flags
mul (unsigned)	reg8	%al*reg8 → %ax	CF,OF
	reg16	%ax*reg16 → %dx.%ax	
	reg32	%al*reg32 → %eax	
imul (signed)	reg8	%al*reg8 → %ax	
	reg16	%ax*reg16 → %dx.%ax	
	reg32	%al*reg32 → %eax	
div (unsigned)	reg8	%ax/reg8 → %al ; %ax mod reg8 → %ah	OF,SF, ZF,AF, CF,PF
	reg16	%ax/reg16→%al ; %ax mod %reg16→%ah	
	reg32	%eax/reg32→%eax ; %dx.%ax mod reg32→%dx	
idiv (signed)	reg8	%ax/reg8→%al ; %ax mod reg8→%ah	
	reg16	%ax/reg16→%al ; %ax mod reg16→%ah	
	reg32	%ax/reg32→%al	



$$-15*(8/2+3)$$

```
.globl main
```

```
main:
```

```
    movl    $8, %eax
```

```
    movl    $0, %edx
```

```
    movl    $2, %ebx
```

```
    divl    %ebx           # %eax = %edx.%eax / %ebx
```

```
    addl    $3, %eax       # %eax = %eax+3
```

```
    movl    $-15, %ebx
```

```
    imull   %ebx           # %eax = -15*%eax
```

```
    ret
```



Jump Operators

Mnemonic	Operand	Operation	Notes
jmp	lbl	jump to lbl	-
ja/jne	lbl	Jump if above / not below or equal	unsigned operands
jae/jnb	lbl	Jump if above or equal / not below	
jbe/jna	lbl	Jump if below or equal / not above	
jb/jnae	lbl	Jump if below / not above or equal	
jg/jnle	lbl	Jump if greater / not less or equal	signed operands
jge/jnl	lbl	Jump if greater or equal / not less	
jle/jng	lbl	Jump if less or equal / not greater	
j1/jnge	lbl	Jump if less / not greater or equal	
je/jz	lbl	Jump if equal / zero (ZF=1)	equality testing
jne/jnz	lbl	Jump if not equal / not zero (ZF=0)	
jc	lbl	Jump if (CF=1)	-
jnc	lbl	Jump if (CF=0)	
js	lbl	Jump if (SF=1)	
jns	lbl	Jump if (SF=0)	



IF ... THEN ... ELSE ...

```
.globl main
```

```
main:
```

```
    movl    $8, %ebx  
    cmpl    %eax, %ebx  
    jle     L0  
    incl    %ebx  
    ret
```

```
# Compare %eax, %ebx  
# If %eax ≤ %ebx go to L0  
# Increment %ebx (then)
```

```
L0:    decl    %ebx  
    ret
```

```
# Decrement %ebx (else)
```



Loops

Mnemonic	Operand	Operation
loop	lbl	%cx-1→%cx ; if (%cx!=0) jump to lbl
loope	lbl	%cx-1→%cx ; if (%cx<>0) and (ZF=1) jump to lbl
loopne	lbl	%cx-1→%cx ; if (%cx<>0) and (ZF=0) jump to lbl
loopz	lbl	%cx-1→%cx ; if (%cx<>0) and (ZF=1) jump to lbl
loopnz	lbl	%cx-1→%cx ; if (%cx<>0) and (ZF=0) jump to lbl

Example:

```
.globl main
main:

    movl    $3,    %ecx
L0:  addl    $5,    %eax
    loop    L0

    ret
```



Floating-point Unit Instruction Set



Moving Data in FPU

Mnemonic	Operand	Operation
finit	-	FPU Initialization
fincstp	-	Increment the FPU stack pointer
fdecstp	-	Decrement the FPU stack pointer
ffree	st(i)	Free the content of st(i)
fldz	-	Load zero in st(0)
fld1	-	Load one in st(0)
fldpi	-	Load π in st(0)
fld	?	Load a float in st(0)
fild	?	Load an int in st(0)
fst	?	Write a float in main memory
fstp	?	Write a float in main memory and pop
fxch	?	Exchange two registers content



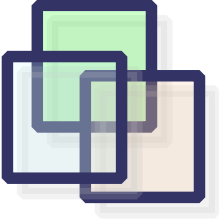
FPU Arithmetic

Mnemonic	Operand	Operation
fadd	-	$st(0) + st(1) \rightarrow st(0)$
	mem/st(i)	$st(0) + mem/st(i) \rightarrow st(0)$
	mem/st(i) mem/st(j)	$st(0) + mem/st(i) \rightarrow st(0)$
fsub	-	Similar to fadd but for substraction
fmul/fdiv	-	Similar to fadd but for multiplication/division
fchs	-	Change sign
fabs	-	Absolute value
fsqr	-	Square
fsqrt	-	Square root
fcos	-	Cosine
fyl2x	-	$y * \log_2(x)$
frndint	-	Round to integer value



FPU Conditionals

Mnemonic	Operand	Operand	Operation
fcom	mem/st(i)	mem/st(j)	Compare two operands, store result in STW
fcomp	mem/st(i)	mem/st(j)	Compare two operands, store result in STW and pop
fcomi	mem/st(i)	mem/st(j)	Compare two int operands, store result in STW
fcomip	mem/st(i)	mem/st(j)	Compare two int operands, store result in STW and pop



Interruptions & System Calls



Interruptions

What does an interruption do:

1. Stop current activity of CPU and save the status;
2. Call a specific subroutine (interrupt handler);
3. Depending on the interruption call (0-255), the interrupt handler load an interrupt vector which jumps to the corresponding subroutine;
4. If several interruptions are occurring at the same time, CPU has a priority order to apply;
5. When the subroutine is finished the CPU restore the CPU status and restart previous execution.



Types of Interrupts

Internal Hardware Interrupts:

Event occurring during the execution of a program
(e.g. division by zero, overflow, general protection error, ...);

External Hardware Interrupts:

Event produced by controllers of external devices
(e.g. PCI/AGP bus, hard-drive, graphic cards, keyboard, ...);

Software Interrupts:

Event produced by programs (mainly by the OS). These interrupts can be produced by using the instruction `int`.



List of Interrupts (Linux)

ID	Message
0x00	Division Error
0x01	Single Step Mode (Debug)
0x02	NMI Interrupt
0x03	Breakpoint
0x04	Overflow
0x05	Bound Range Exceeded
0x06	Invalid Opcode
0x07	Coprocessor Not Available
0x08	Double Exception
0x09	Coprocessor Segment Overrun
0x0a	Invalid Task State Segment
0x0b	Segment Not Present
0x0c	Stack Fault
0x0d	General Protection
0x0e	Page Fault
0x0f	Reserved
0x10	Coprocessor Error
0x11-0x1f	Reserved
0x12-0xffffffff	Coprocessor Error



System Calls (Linux)

A system call is a software interrupt tight to a specific subroutine of the OS. (e.g. get a char from keyboard, print on stdout, ...).

- **Arguments:**

- `%eax`: Syscall ID
- `%ebx,%ecx,%edx,%esi,%edi`: Syscall arguments

- **Calling a syscall:**

- `int $0x80`

- **Example:**

```
.globl main
main:
    movl    $1,%eax    # Interruption ID
    int     $0x80      # Calling the OS
    ret
```




Few System Calls (Linux)

%eax	Name	%ebx	%ecx	%edx	%esi	%edi
1	sys_exit	int	-	-	-	-
2	sys_fork	struct pt_regs	-	-	-	-
3	sys_read	unsigned int	char*	size_t	-	-
4	sys_write	unsigned int	const char*	size_t	-	-
5	sys_open	const char*	int	int	-	-
6	sys_close	unsigned int	-	-	-	-
7	sys_waitpid	pid_t	unsigned int	int	-	-
8	sys_create	const char*	int	-	-	-
9	sys_link	const char*	const char*	-	-	-
10	sys_unlink	const char*	-	-	-	-
11	sys_execve	struct pt_regs	-	-	-	-
12	sys_chdir	const char*	-	-	-	-
13	sys_time	int*	-	-	-	-
14	sys_mknod	const char*	mode_t	dev_t	-	-
15	sys_chmod	const char*	mode_t	-	-	-
...

See in linux-2.6.x.y/arch/i386/kernel/syscall_table.S



Hello World !

```
.data                                # Data section
msg:
    .ascii "Hello World!\n"         # String
    len = . -msg                    # String length

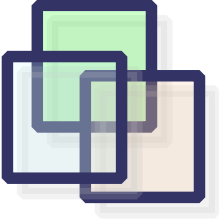
.text                                # Text section
    .globl main                      # Export entry point to ELF linker

main:
# Write the string to stdout
    movl    $len, %edx               # 3rd argument: string length
    movl    $msg, %ecx               # 2nd argument: pointer to string
    movl    $1, %ebx                 # 1st argument: file handler (stdout)
    movl    $4, %eax                 # System call number (sys_write)
    int     $0x80                    # Kernel call
# and exit
    movl    $0, %ebx                 # 1st argument: exit code
    movl    $1, %eax                 # System call number (sys_exit)
    int     $0x80                    # Kernel call
```

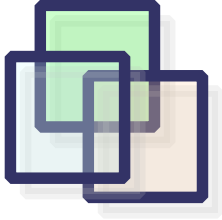


Other Instructions

- **hlt:**
Stop the CPU until the next interruption occurs
- **nop:**
No Operation. Used to avoid to stall the pipeline (waiting for data to evaluate a test)
- **wait:**
Deprecated instruction used to wait for the results coming from arithmetic co-processor



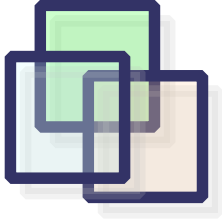
Advanced Instruction Sets



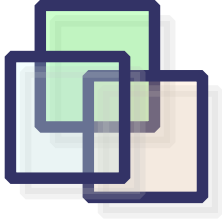
MMX Instructions



SSE Instructions



SSE2 Instructions



SSE3 Instructions



3DNow! Instructions



Next Time

**From C to Assembly
(and back...)**